

"S" +



=

Spark-Bench

Simulate, Test, Compare, Exercise, and Yes, Even
Benchmark!

Emily May Curtin

IBM Watson Data Platform
Atlanta, GA



@emilymaycurtin



@ecurtin

#EUeco8

#sparkbench

Who Am I

- Artist, former documentary film editor
- 2017 Spark Summit East – Spark + Parquet
- Tabs, not spaces (at least for Scala)(where 1 tab = 2 spaces)
- Software Engineer, IBM Watson Data Platform
- **Current lead developer of SparkTC/spark-bench**

Atlanta!!!



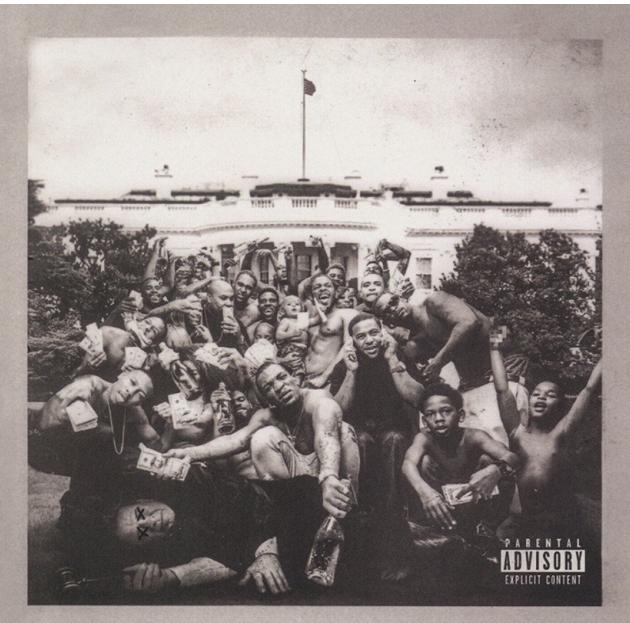
Spark-Bench

- History & Rewrite Motivation
- How Spark-Bench Works
- Detailed Example: Parquet vs. CSV Benchmarking
- More Examples
 - Cluster Hammer
 - Multiple Notebook Users
 - Varying over Spark parameters

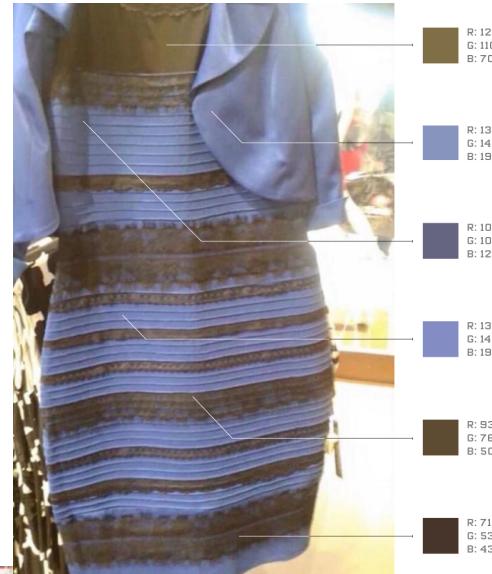
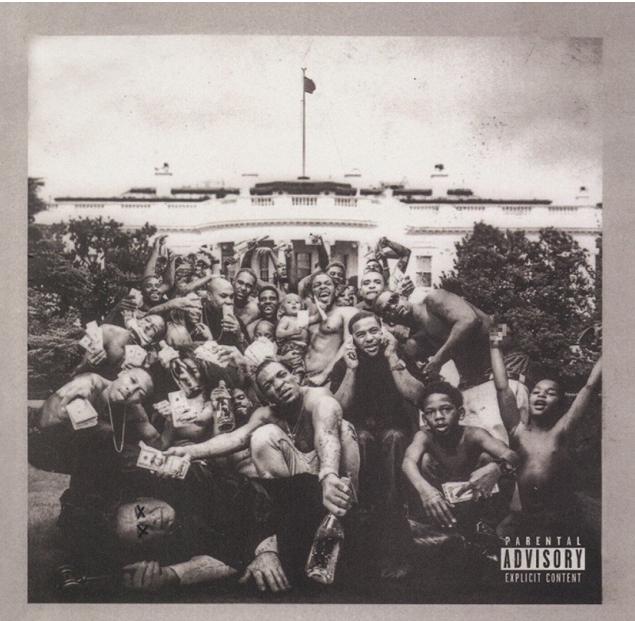
2015 A.D.



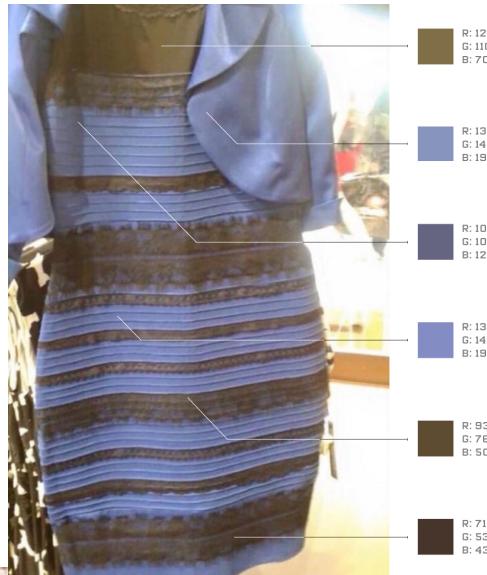
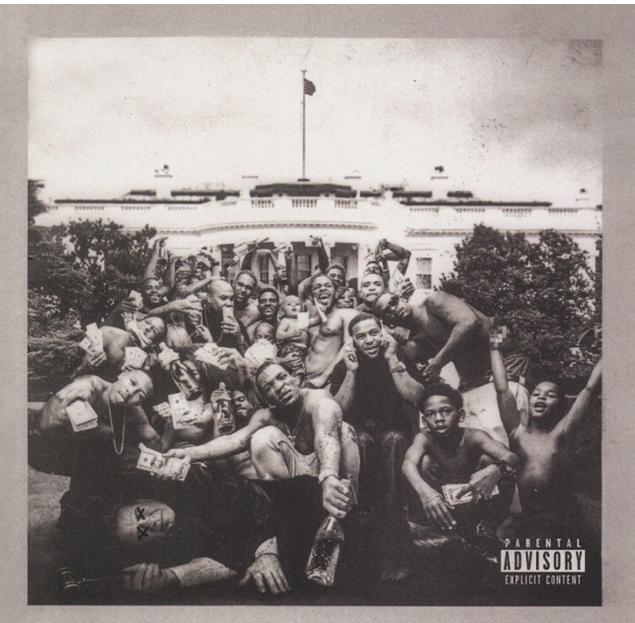
2015 A.D.



2015 A.D.



2015 A.D.



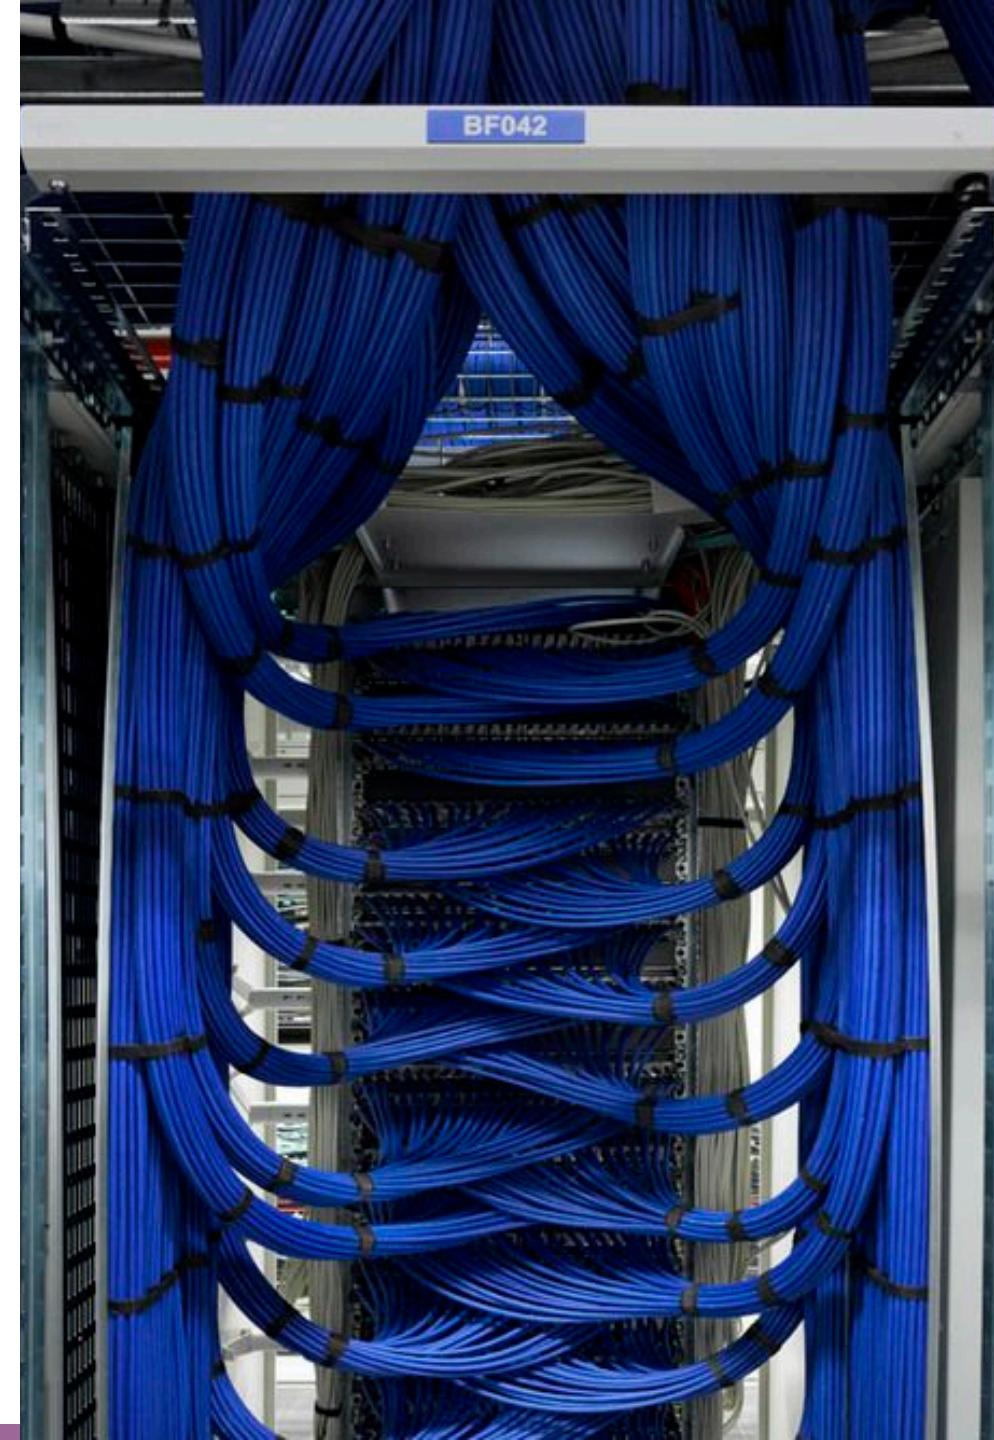
Spark Benchmarking circa 2015 A.D.



Spark-Bench 2015

"SparkBench: A Comprehensive Benchmarking Suite For In Memory Data Analytics Platform Spark"

Min Li, Jian Tan, Yandong Wang, Li Zhang, Valentina Salapura IBM TJ Watson Research Center - APC 2015



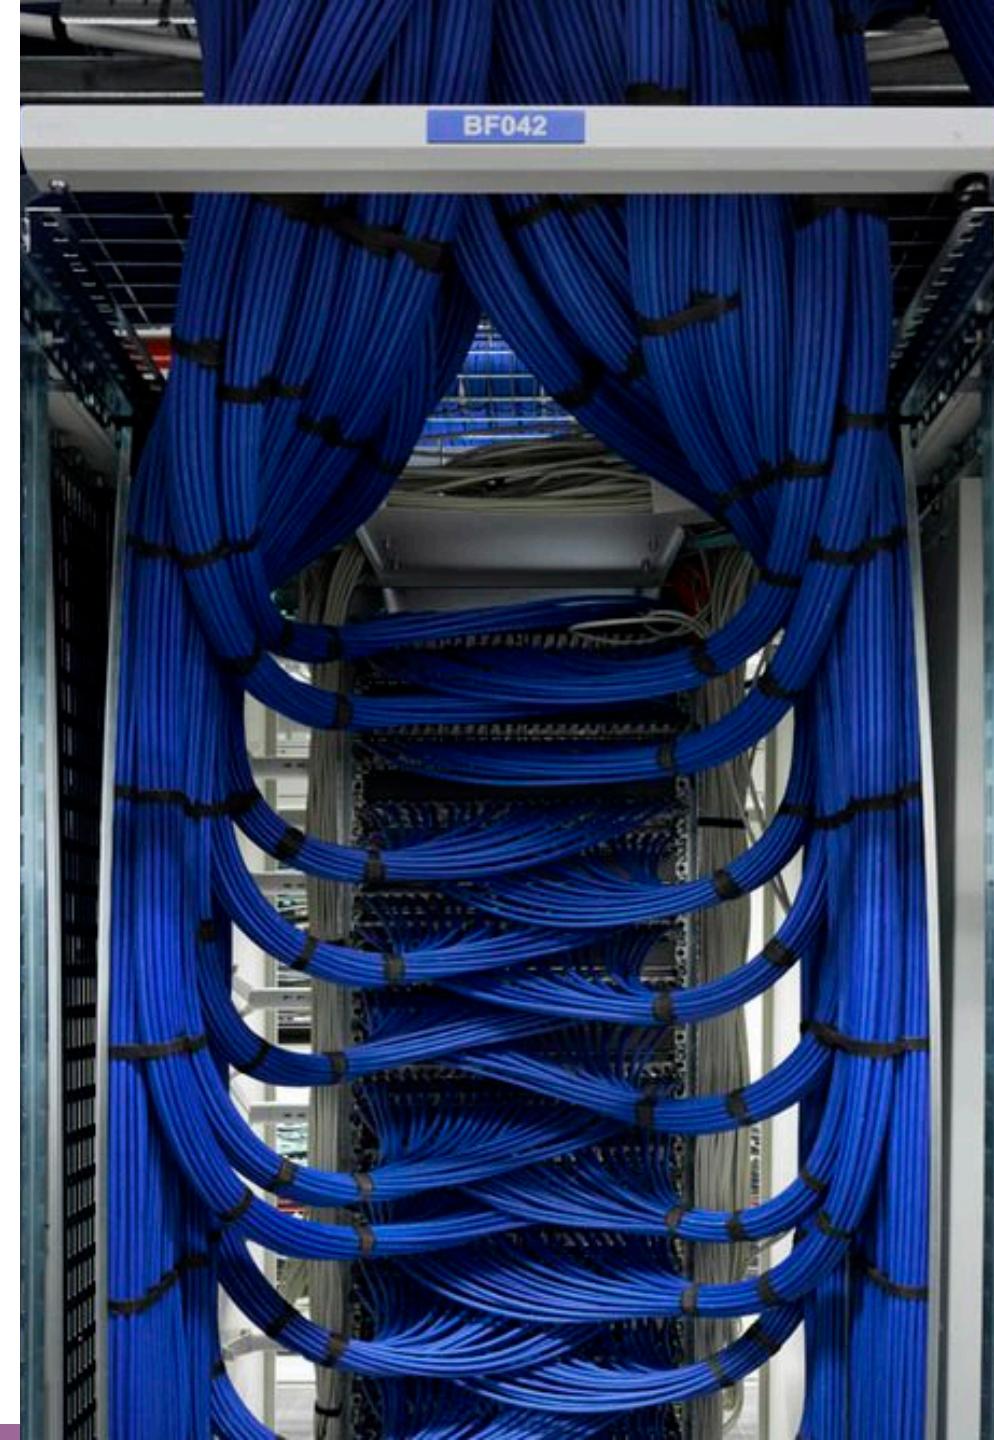
Spark-Bench 2015

"SparkBench: A Comprehensive Benchmarking Suite For In Memory Data Analytics Platform Spark"

Min Li, Jian Tan, Yandong Wang, Li Zhang, Valentina Salapura IBM TJ Watson Research Center - APC 2015



#EUec08 #sparkbench

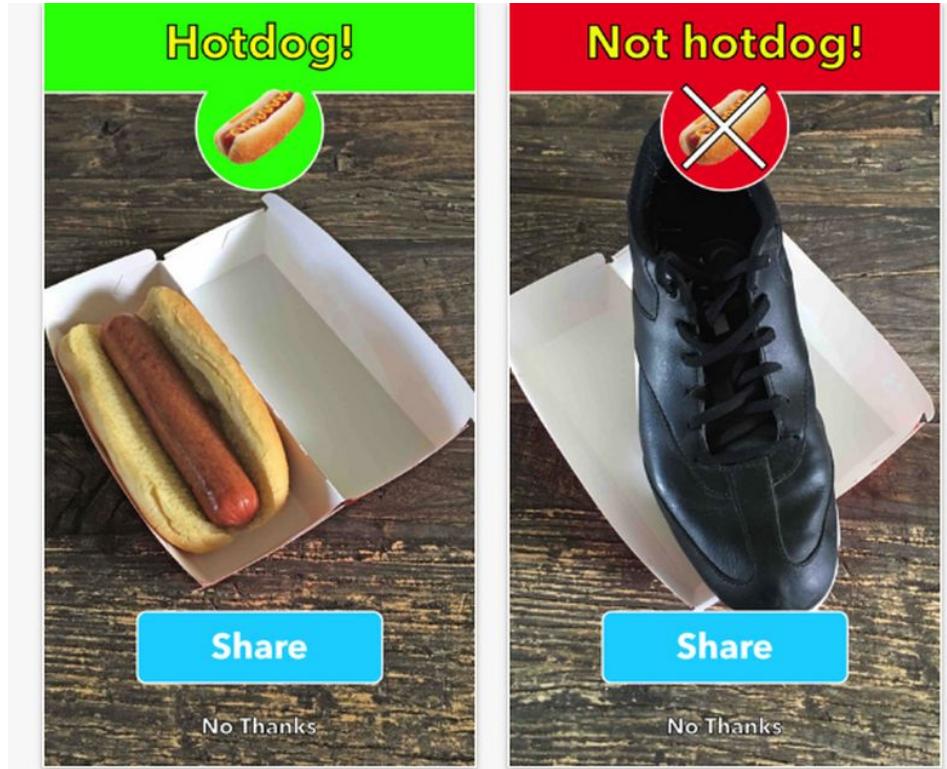


2017 A.D.

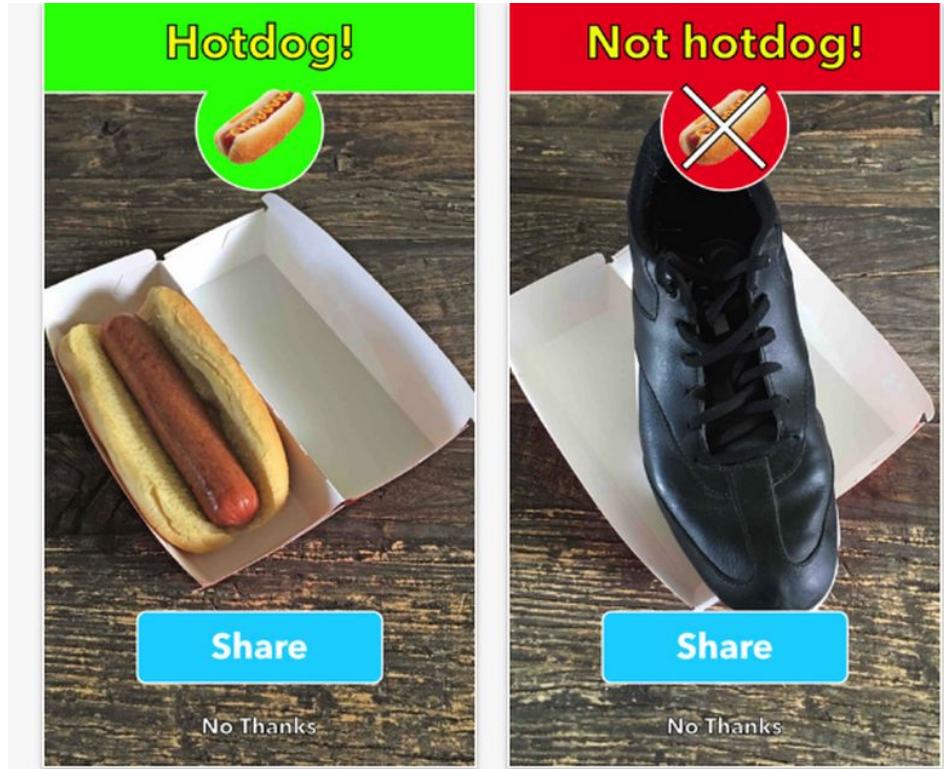
2017 A.D.



2017 A.D.



2017 A.D.



Watson Data Platform

Data Science
Experience



IBM Analytics
Engine (*Beta*)



Watson Machine
Learning

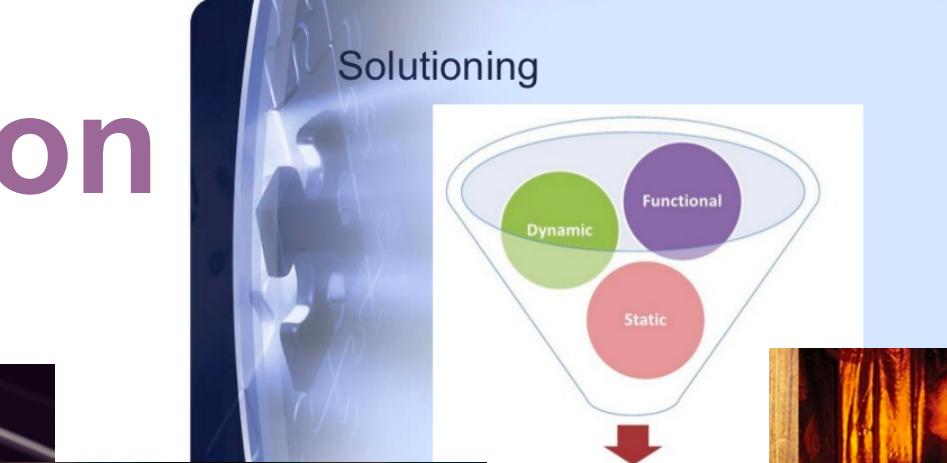
AND MUCH,
MUCH MORE

**As a researcher, I want to run a set
of standard benchmarks with
generated data against one
instance of Spark, one time.**

--The user story that SparkTC/spark-bench version 1.0 fulfilled

Good start, let's keep going!

The Vision



**As a Spark developer, I want to
run regression tests with my
changes in order to provide some
authoritative performance stats
with my PR.**

**As a Spark developer, I want to
see regression tests run against
Spark nightly builds.**

**As a Parquet developer, I want to
compare performance between
Spark built with several different
versions of Parquet.**

As a Spark developer, I want to simulate multiple notebook users hitting the same cluster so that I can test my changes to the dynamic allocation algorithm.

**As an enterprise user of Spark, I
want to make my daily batch
ingest job run faster by tuning
parameters.**

**As a DevOps professional, I want
to see how performance changes
for our daily jobs as I add more
nodes to our prod cluster.**

As a data scientist, I want to run several machine learning algorithms over the same set of data in order to compare results.

As anybody who has ever had to make a graph about Spark, I want to get numbers easily and quickly so I can focus on making *beautiful graphs*.



How It Works

#EUec08 #sparkbench

Structural Details

	Old Version	✨ New Version ✨
Project Design	Series of shell scripts calling individually built jars	Scala project built using SBT
Build/Release	Manually/Manually	SBT, TravisCI, auto-release to Github Releases on PR merge
Configuration	Scattered in shell script variables	Centralized in one config file with many new capabilities
Parallelism	None	Three different levels!
Custom Workloads	Requires writing a new subproject with all accompanying bash	Implement our abstract class, Bring Your Own Jar

The Config File (a first glance)

```
spark-bench = {  
    spark-submit-config = [{  
        workload-suites = [  
            {  
                descr = "One run of SparkPi and that's it!"  
                benchmark-output = "console"  
                workloads = [  
                    {  
                        name = "sparkpi"  
                        slices = 10  
                    }  
                ]  
            }  
        ]  
    }]  
}
```

#EUec08 #sparkbench

Workloads

- Atomic unit of spark-bench
- Optionally read from disk
- Optionally write results to disk
- All stages of the workload are timed
- Infinitely composable

Workloads

- Machine Learning
 - Kmeans
 - Logistic Regression
 - Other workloads ported from old version
- SQL queries
- Oddballs
 - SparkPi
 - Sleep
 - Cache Test

Workloads: Data Generators

- Graph data generator
- Kmeans data generator
- Linear Regression data generator
- More coming!

What About Other Workloads?

1. Active development, more coming!
2. Contributions welcome ☺

3. Bring Your Own Workload!

```
{  
    name = "custom"  
    class = "com.seriouscompany.HelloString"  
    str = ["Hello", "Hi"]  
}
```

Custom Workloads

```
{  
    name = "custom"  
    class = "com.seriouscompany.OurGiganticIngestDailyIngestJob"  
    date = "Tues Oct 26 18:58:25 EDT 2017"  
    source = "s3://buncha-junk/2017-10-26/"  
}
```

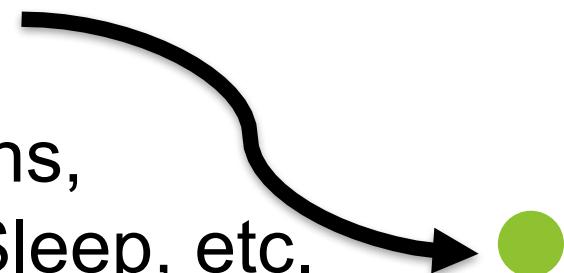
```
{  
    name = "custom"  
    class = "com.example.WordGenerator"  
    output = "console"  
    rows = 10  
    cols = 3  
    word = "Cool stuff!!"  
}
```

← See our documentation on
custom workloads!



Workload

Could be Kmeans,
SparkPi, SQL, Sleep, etc.



```
{  
  name = "sparkpi"  
  slices = 500000000  
}
```

SparkPi

Borrowing from the classic Spark example, this workload computes an approximation of pi. From the Spark examples page:
<https://spark.apache.org/examples.html>

Spark can also be used for compute-intensive tasks. This code estimates π by "throwing darts" at a circle. We pick random points in the unit square ((0, 0) to (1,1)) and see how many fall in the unit circle. The fraction should be $\pi / 4$, so we use this to get our estimate.

SparkPi is particularly useful for exercising the computing power of Spark without the consideration of heavy I/O from data-reliant workloads.

The timing result of SparkPi will include the estimate of Pi that was generated.

Parameters

Name	Required	Default	Description
name	yes	-	"sparkpi"
slices	no	2	Number of partitions that will be spawned

Examples

{

#EUec08 #sparkbench

Workload

Could be Kmeans,
SparkPi, SQL, Sleep, etc.

```
{  
    name = "sparkpi"  
    slices = 50000000  
}
```

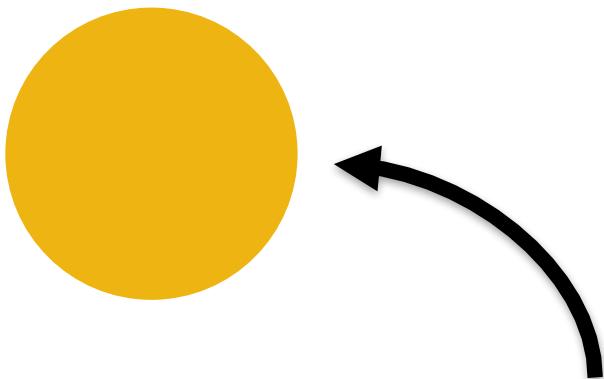
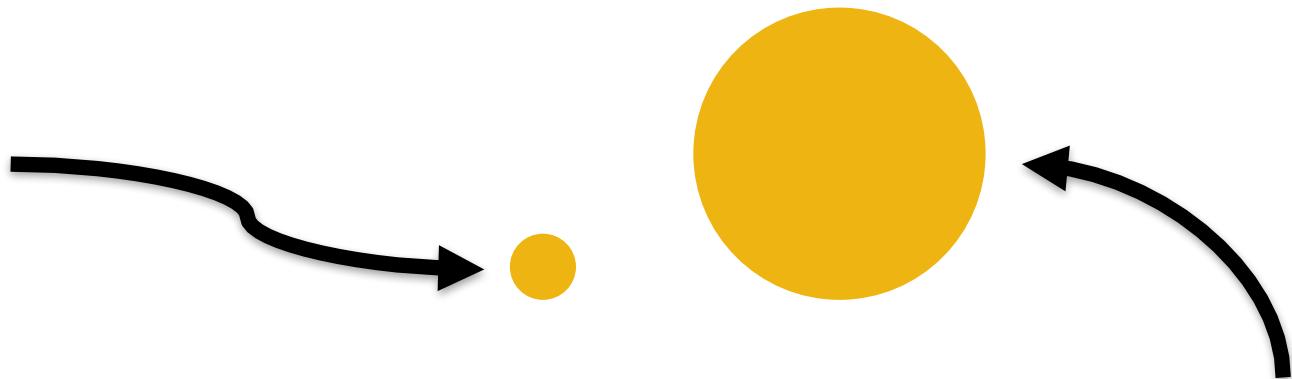


Data Generation Workload

```
{  
    name = "data-generation-kmeans"  
    rows = 1000  
    cols = 24  
    output = "/tmp/kmeans-data.parquet"  
    k = 45  
}
```

Data Generation, Small Set

```
{  
    name = "data-generation-kmeans"  
    rows = 100  
    cols = 24  
    output = "/tmp/small-data.parquet"  
}
```



Data Generation, Large Set

```
{  
    name = "data-generation-kmeans"  
    rows = 100000000  
    cols = 24  
    output = "/tmp/large-data.parquet"  
}
```

TIME



Workloads running
serially

TIME

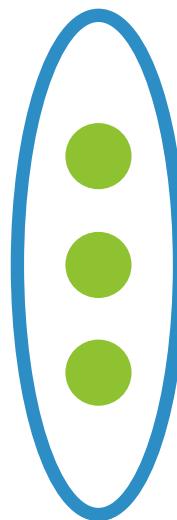


Workloads running
in parallel

Workload Suites

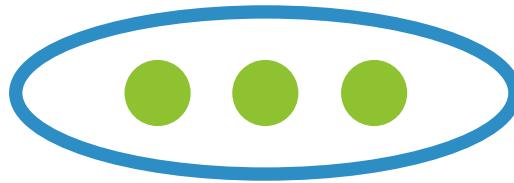
- Logical group of workloads
- Run workloads serially or in parallel
- Control the repetition of workloads
- Control the output of benchmark results
- Can be composed with other workload suites to run serially or in parallel

One workload suite
with three workloads
running **serially**



```
workload-suites = [ {  
    descr = "Lol"  
    parallel = false  
    repeat = 1  
    benchmark-output = "console"  
    workloads = [  
        {  
            // Three workloads here  
        }  
    ]  
} ]
```

One workload suite
with three workloads
running in parallel



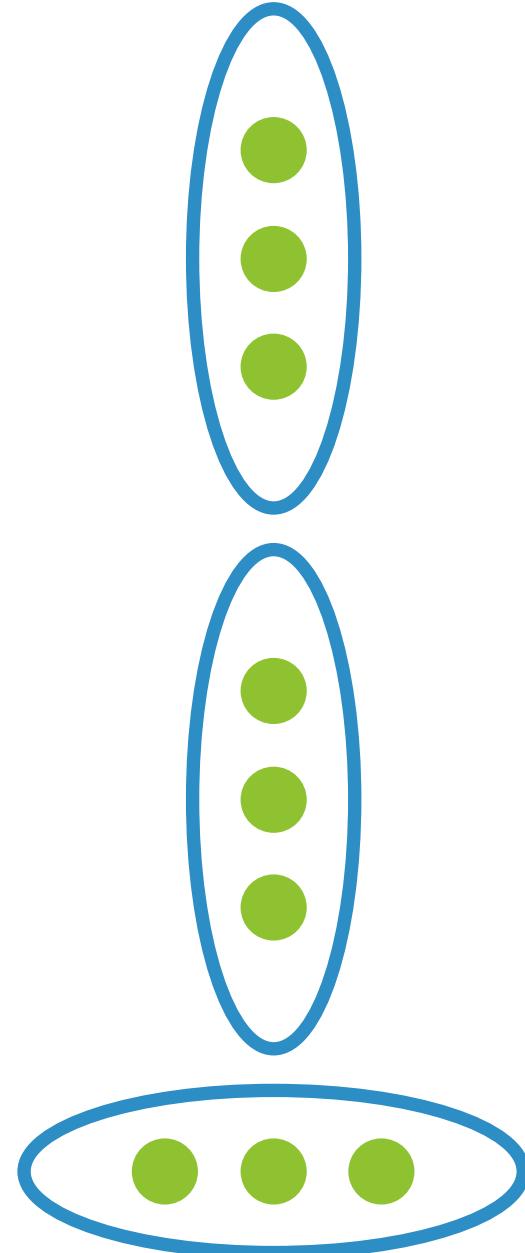
```
workload-suites = [ {  
    descr = "Lol"  
    parallel = true  
    repeat = 1  
    benchmark-output = "console"  
    workloads = [  
        {  
            // Three workloads here  
        }  
    ]  
} ]
```

Three workload suites running **serially**.

Two have workloads running serially,
one has workloads running in parallel.

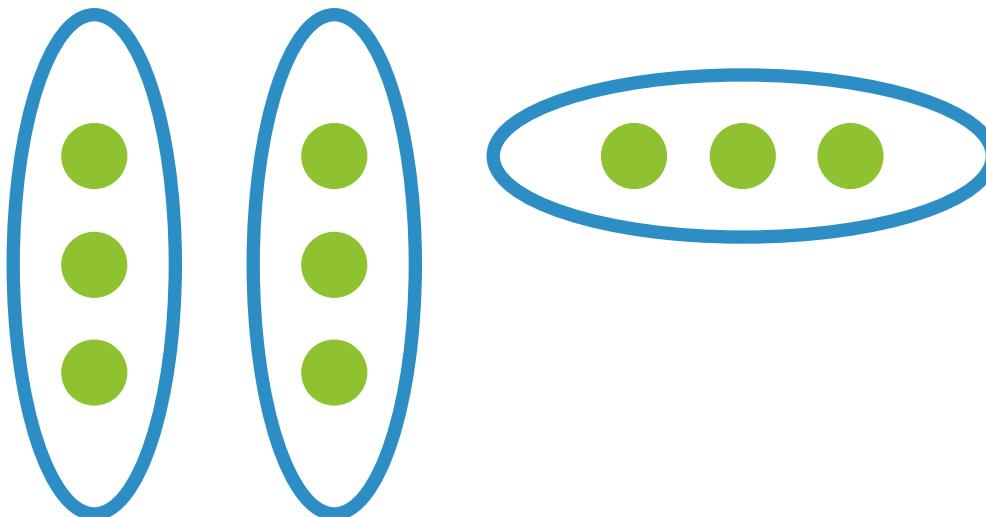
(config file is abbreviated for space)

```
suites-parallel = false
workload-suites = [
    {
        descr = "One"
        parallel = false
    },
    {
        descr = "Two"
        parallel = false
    },
    {
        descr = "Three"
        parallel = true
    }
]
```



Three workload suites running **in parallel**.

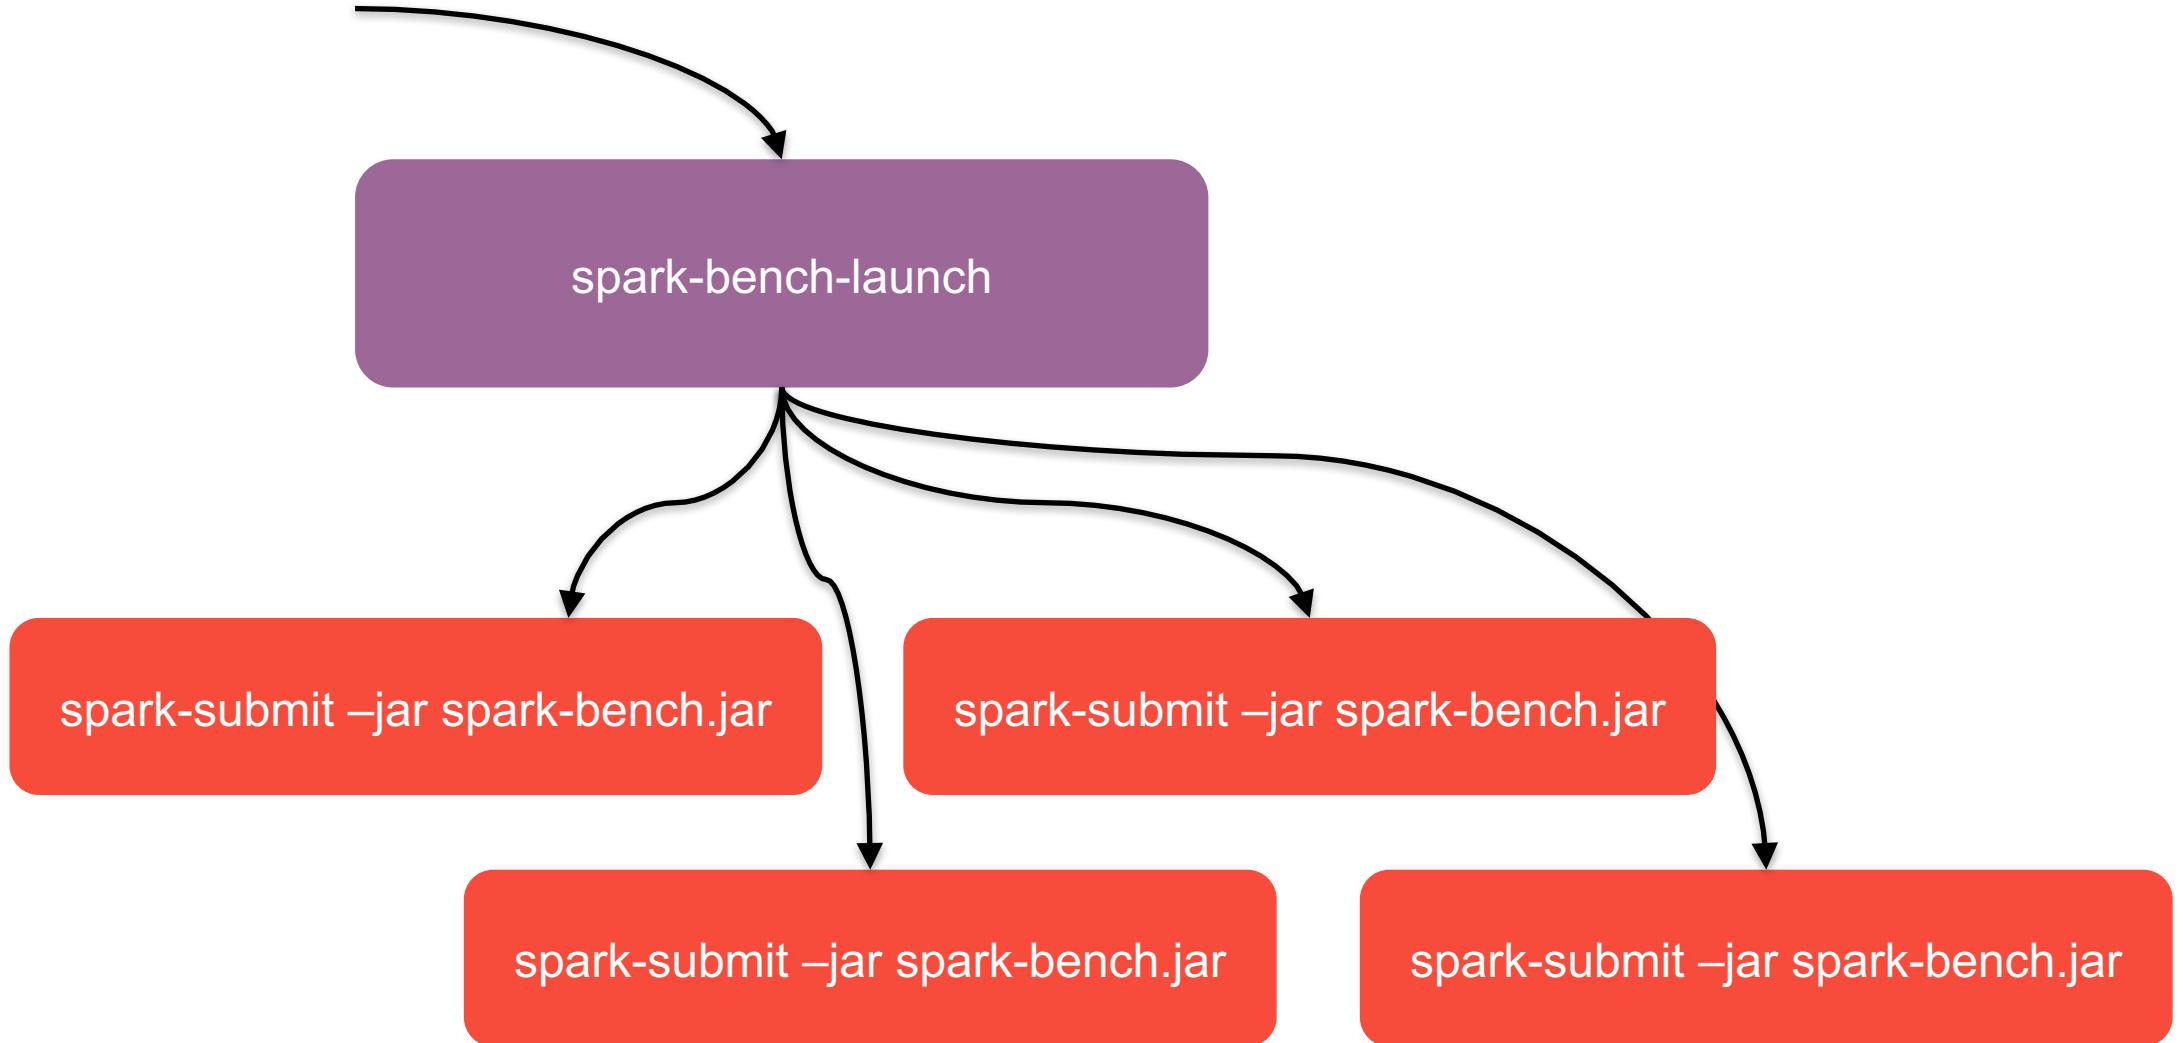
Two have workloads running **serially**,
one has workloads running **in parallel**.



(config file is abbreviated for space)

```
suites-parallel = true
workload-suites = [
  {
    descr = "One"
    parallel = false
  },
  {
    descr = "Two"
    parallel = false
  },
  {
    descr = "Three"
    parallel = true
  }
]
```

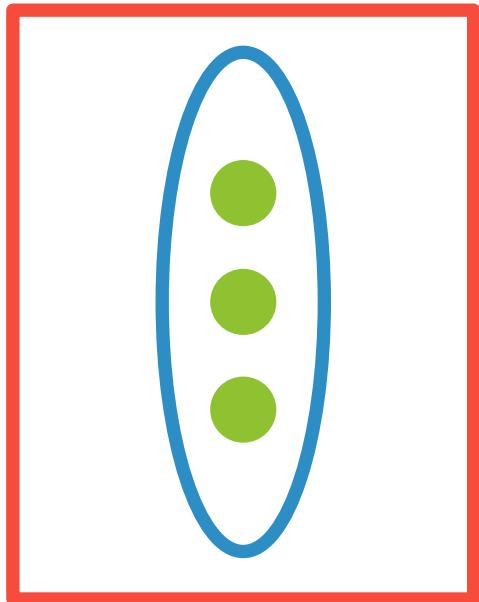
```
./bin/spark-bench.sh my-amazing-config-file.conf
```



Spark-Submit-Config

- Control any parameter present in a spark-submit script
- Produce and launch multiple spark-submit scripts
- Vary over parameters like executor-mem
- Run against different clusters or builds of Spark
- Can run serially or in parallel

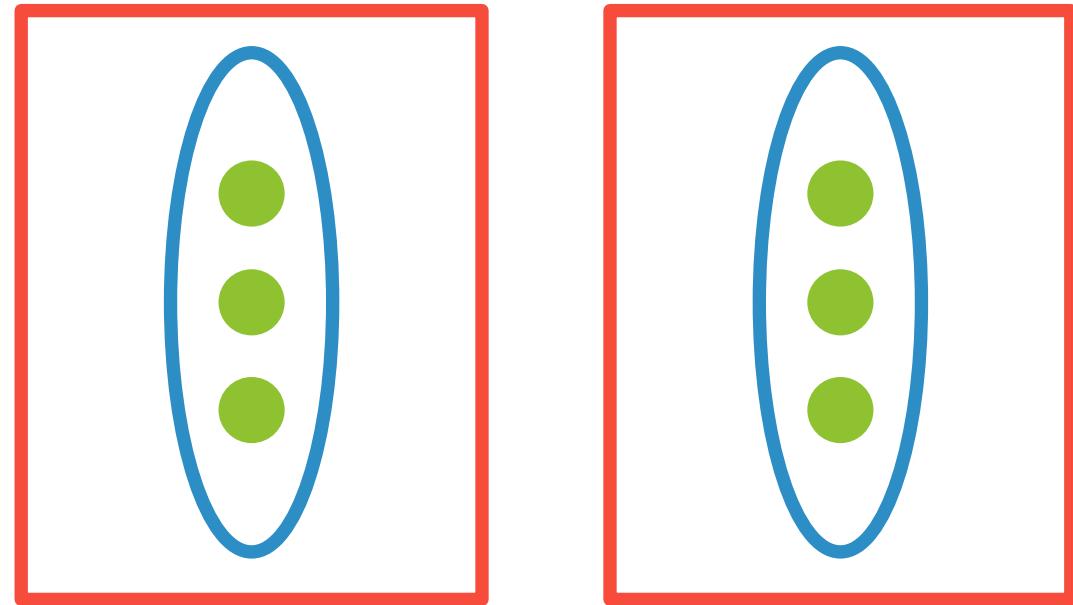
One spark-submit
with one workload
suite with three
workloads running
serially



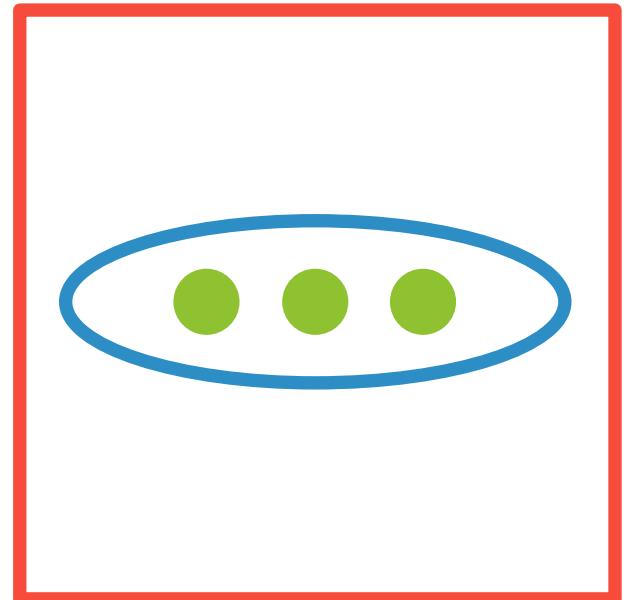
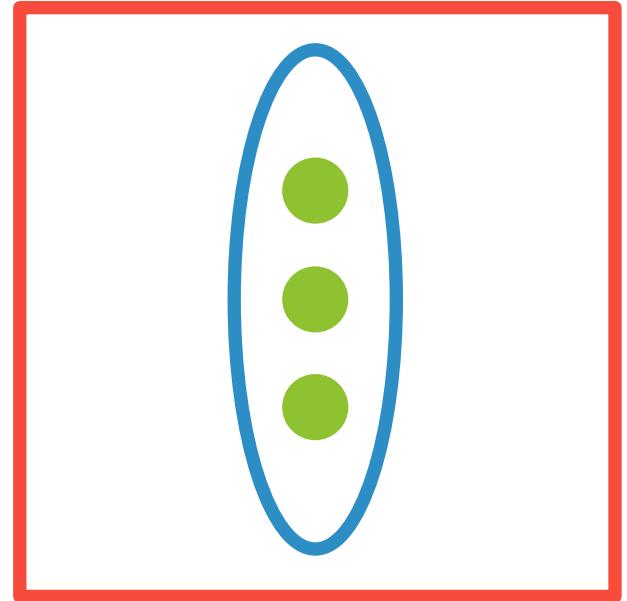
(config file is abbreviated for space)

```
spark-bench = {  
    spark-submit-config = [ {  
        spark-args = {  
            master = "yarn"  
        }  
        workload-suites = [ {  
            descr = "One"  
            benchmark-output = "console"  
            workloads = [  
                // . . .  
            ]  
        } ]  
    } ]  
}
```

Two spark-submits running in parallel, each with one workload suite, each with three workloads running **serially**



Two spark-submits running **serially**,
each with one workload suite,
one with three workloads running **serially**,
one with three workloads running **in parallel**

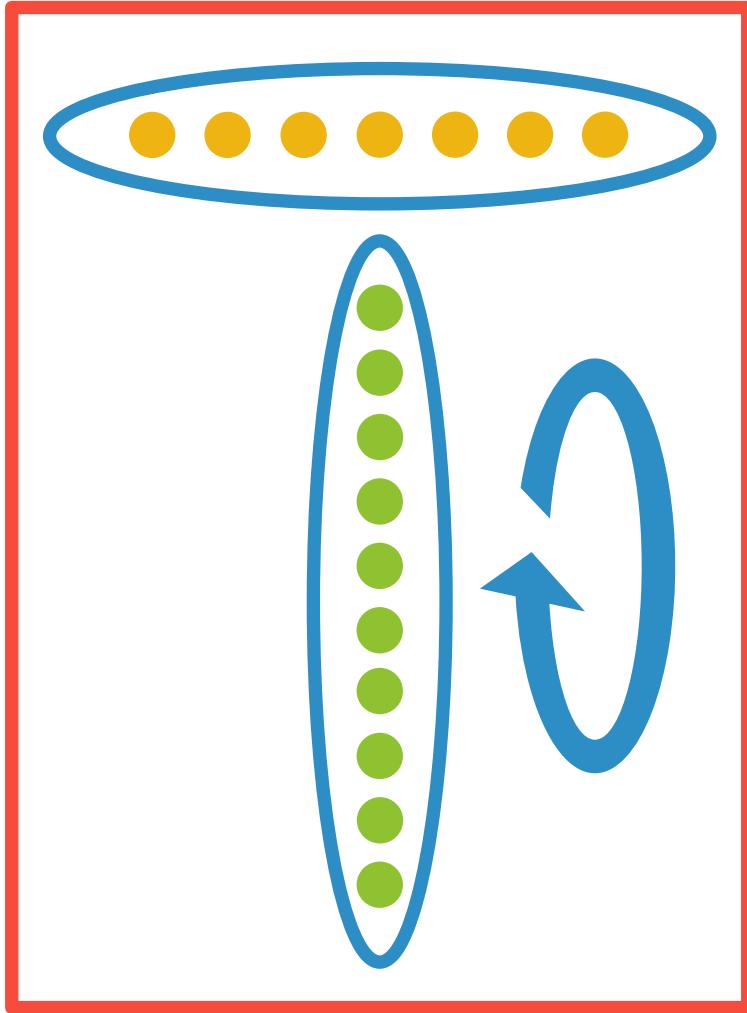


One spark-submit with two workload suites running serially.

The first workload suite generates data **in parallel**.

The second workload suite runs different workloads **serially**.

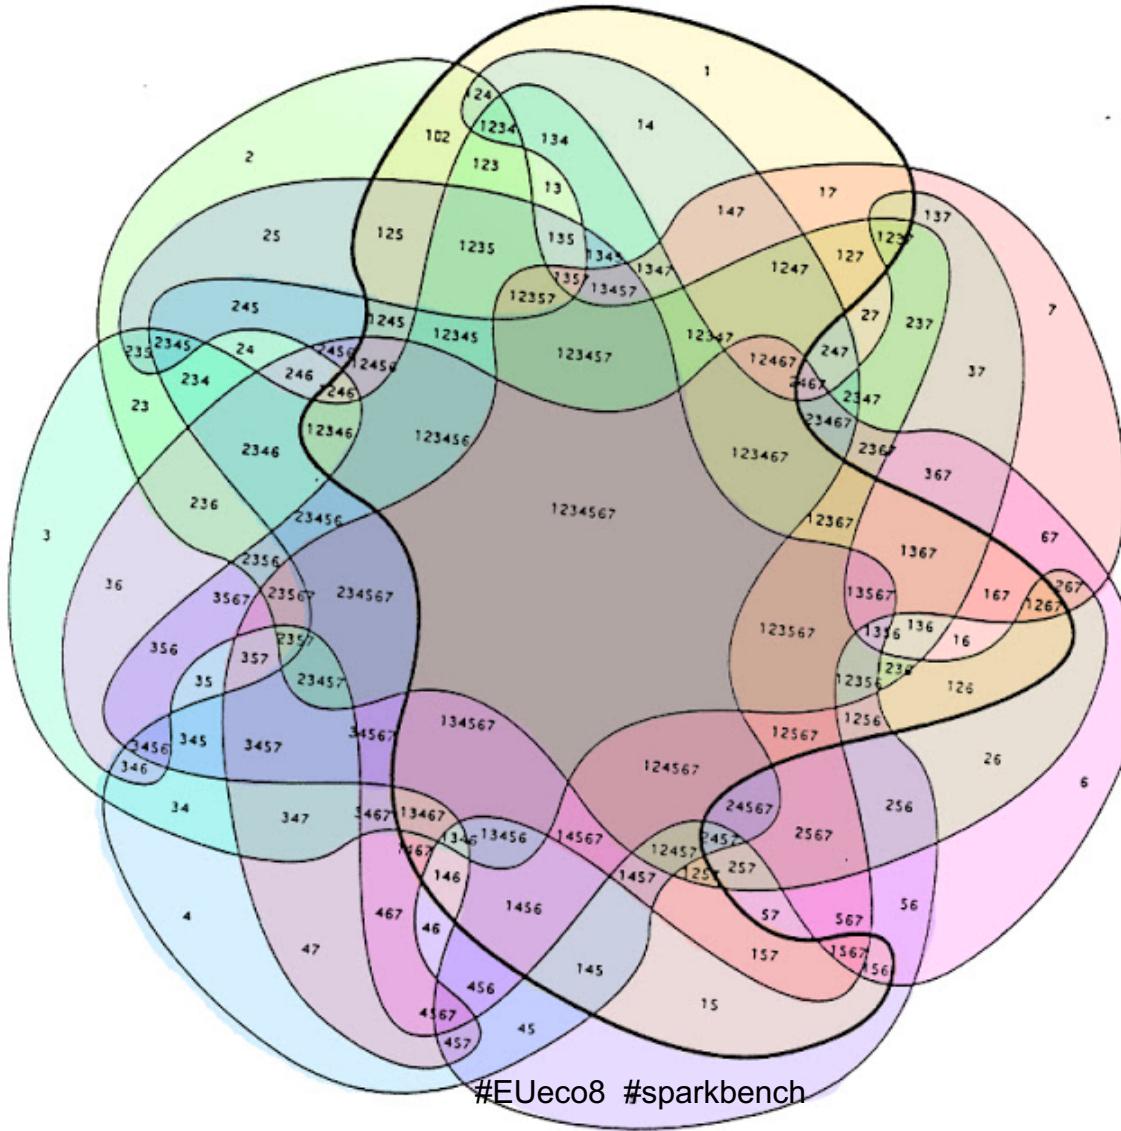
The second workload suite **repeats**.



```
spark-bench = {
    spark-submit-parallel = false
    spark-submit-config = [
        spark-args = { // . . . }
        suites-parallel = false
        workload-suites = [
            descr = "Data Generation"
            parallel = true
            repeat = 1 // generate once and done!
            workloads = { // . . . }
        },
        workload-suites = [
            descr = "Actual workloads that we want to benchmark"
            parallel = false
            repeat = 100 // repeat for statistical validity!
            workloads = { // . . . }
        ]
    ]
}
```

#EUec08 #sparkbench

Beautiful Graphs



Detailed Example: Parquet vs. CSV

Example: Parquet vs. CSV

Goal: benchmark two different SQL queries over the same dataset stored in two different formats

1. Generate a big dataset in CSV format
2. Convert that dataset to Parquet
3. Run first query over both sets
4. Run second query over both sets
5. Make beautiful graphs

```
spark-bench = {
    spark-submit-config = [ {
        spark-home = "/usr/iop/current/spark2-client/"
        spark-args = {
            master = "yarn"
            executor-memory = "28G"
            num-executors = 144
        }
        conf = {
            "spark.dynamicAllocation.enabled" = "false"
            "spark.dynamicAllocation.monitor.enabled" = "false"
            "spark.shuffle.service.enabled" = "true"
        }
        suites-parallel = false
    }
}
```

// ...CONTINUED...

```
workload-suites = [
{
    descr = "Generate a dataset, then transform it to Parquet format"
    benchmark-output = "hdfs://tmp/emily/results-data-gen.csv"
    parallel = false
    workloads = [
        {
            name = "data-generation-kmeans"
            rows = 10000000
            cols = 24
            output = "hdfs://tmp/emily/kmeans-data.csv"
        },
        {
            name = "sql"
            query = "select * from input"
            input = "hdfs://tmp/emily/kmeans-data.csv"
            output = "hdfs://tmp/emily/kmeans-data.parquet"
        }
    ]
},
// ...CONTINUED...
```

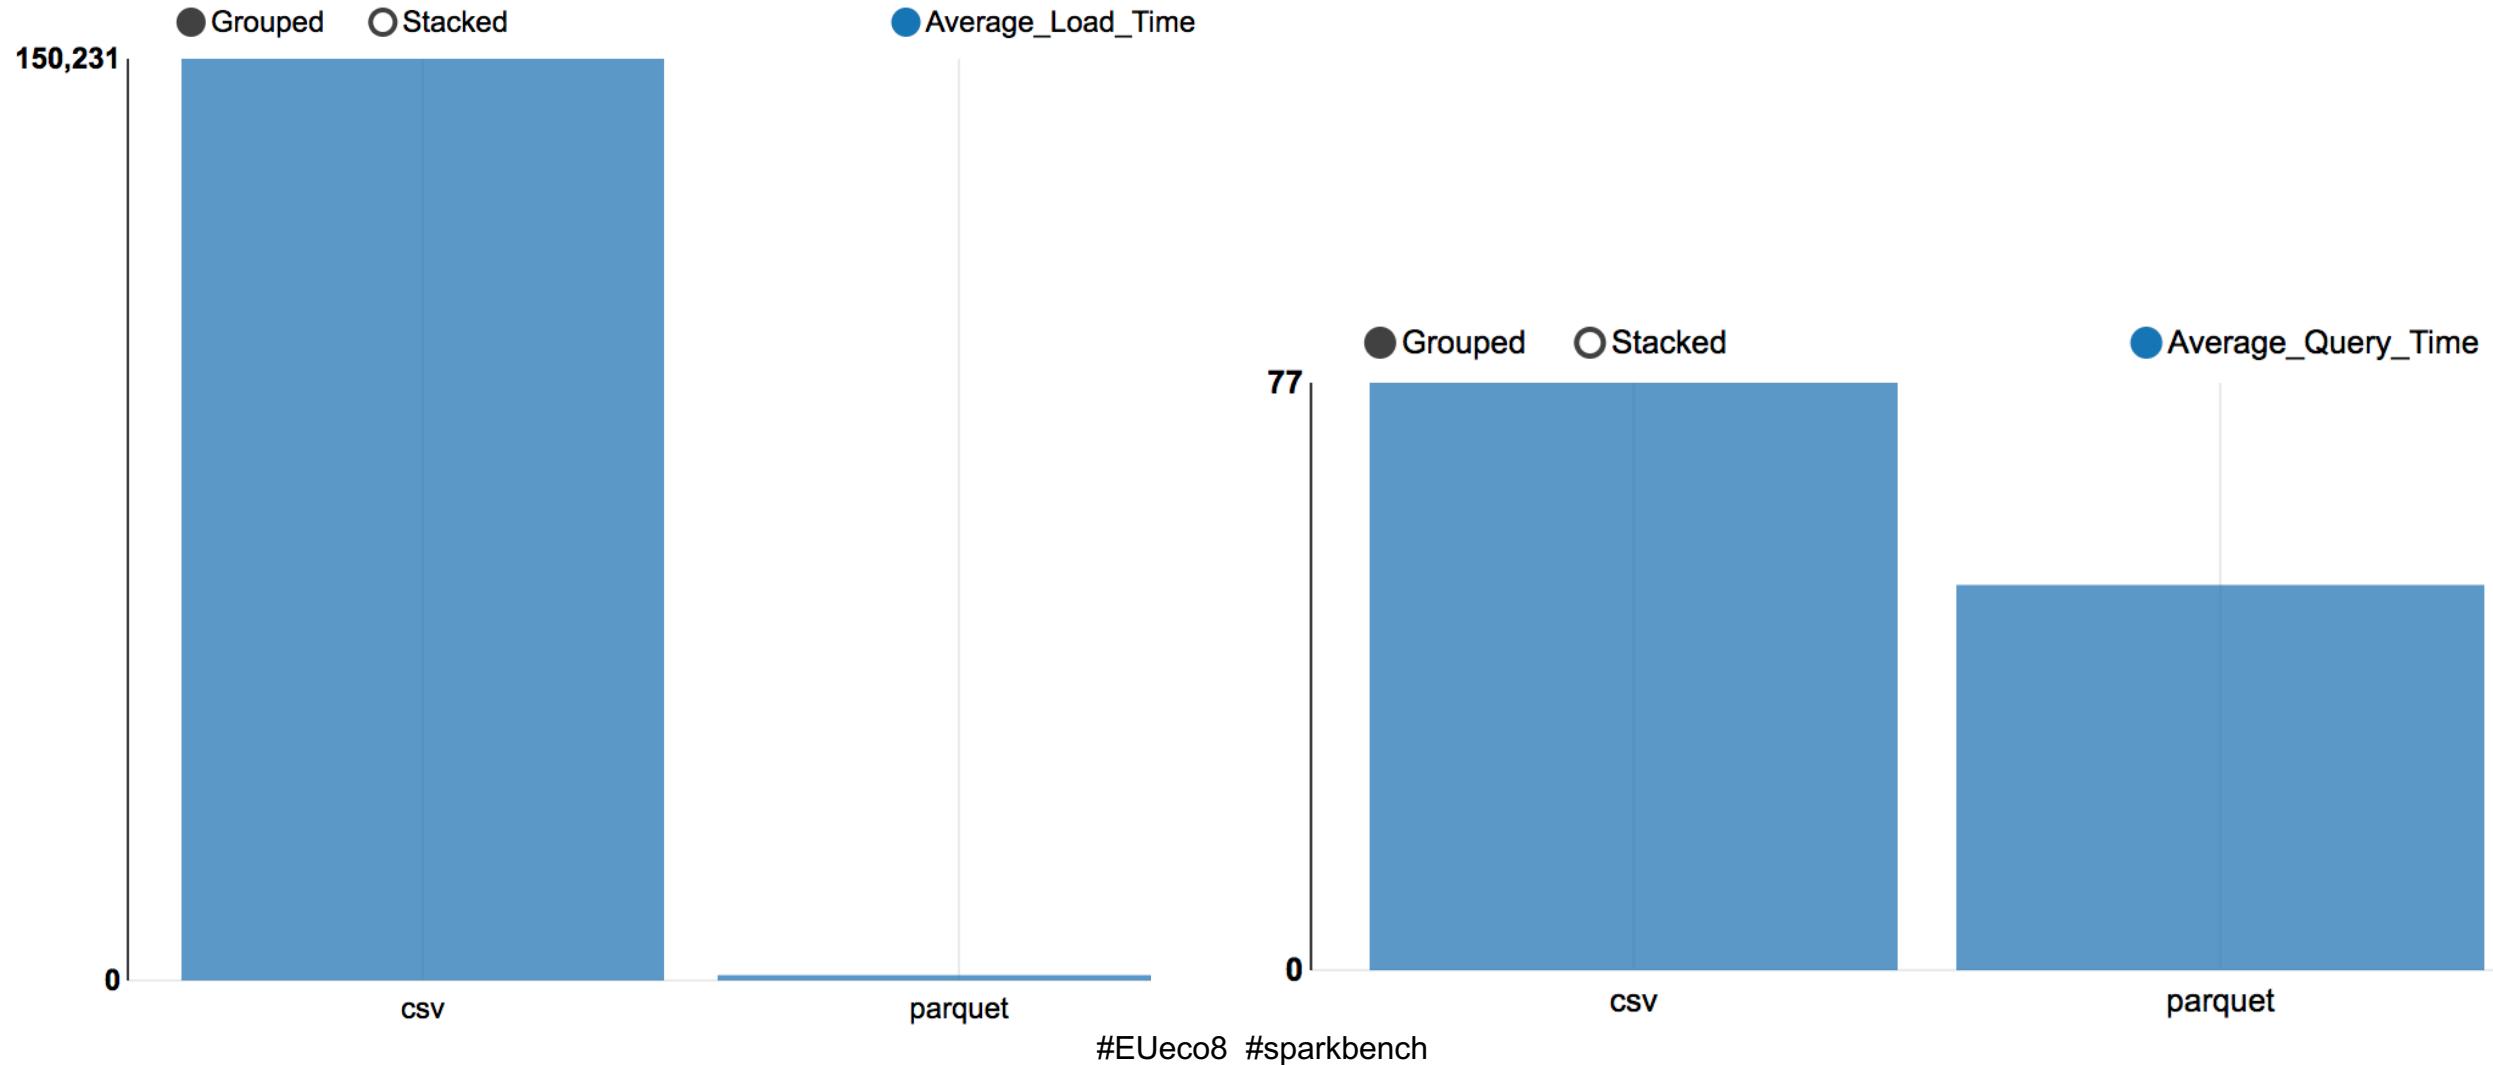
```
{  
    descr = "Run two different SQL over two different formats"  
    benchmark-output = "hdfs://tmp/emily/results-sql.csv"  
    parallel = false  
    repeat = 10  
    workloads = [  
        {  
            name = "sql"  
            input = [  
                "hdfs://tmp/emily/kmeans-data.csv",  
                "hdfs://tmp/emily/kmeans-data.parquet"  
            ]  
            query = [  
                "select * from input",  
                "select `0`, `22` from input where `0` < -0.9"  
            ]  
            cache = false  
        }  
    ]  
}
```



Parquet vs CSV Results

name	timestamp	loadTime	queryTime	total_Runtime	run	cache	queryStr	input	spark.executor.instances	spark.submit.deployMode	spark.master	sp
sql	1.50714E+12	15895947762	8722036	15904669798	0	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	15232039388	41572417	15273611805	0	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	139462342	5652725	145115067	0	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	100552926	9120325	109673251	0	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	15221909508	6183625	15228093133	1	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	15634952815	6820382	15641773197	1	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	93264481	5243927	98508408	1	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	84948895	5573010	90521905	1	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	15270158275	8408445	15278566720	2	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	14924461363	7999643	14932461006	2	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	86821299	13952212	100773511	2	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	100789178	6525153	107314331	2	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	15115682178	7085577	15122767755	3	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	14829435321	6452434	14835887755	3	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	78654779	5291044	83945823	3	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	86702691	5641506	92344197	3	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	14881975292	7346304	14889321596	4	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	14953655628	6749366	14960404994	4	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	102374878	4925237	107300115	4	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	79458610	4697124	84155734	4	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	14908128878	7429576	14915558454	5	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	14922630998	4596120	14927227118	5	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	79144522	4191275	83335797	5	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	86207011	4111788	90318799	5	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	14909544529	5722627	14915267156	6	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	14919548742	5643308	14925192050	6	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.csv		144client	yarn	280
sql	1.50714E+12	85241859	4417683	89659542	6	FALSE	select * from input	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280
sql	1.50714E+12	93379964	4610765	97990729	6	FALSE	select `0`, `22` from input where `0` < -0.9	hdfs://tmp/emily/kmeans-data.parquet		144client	yarn	280

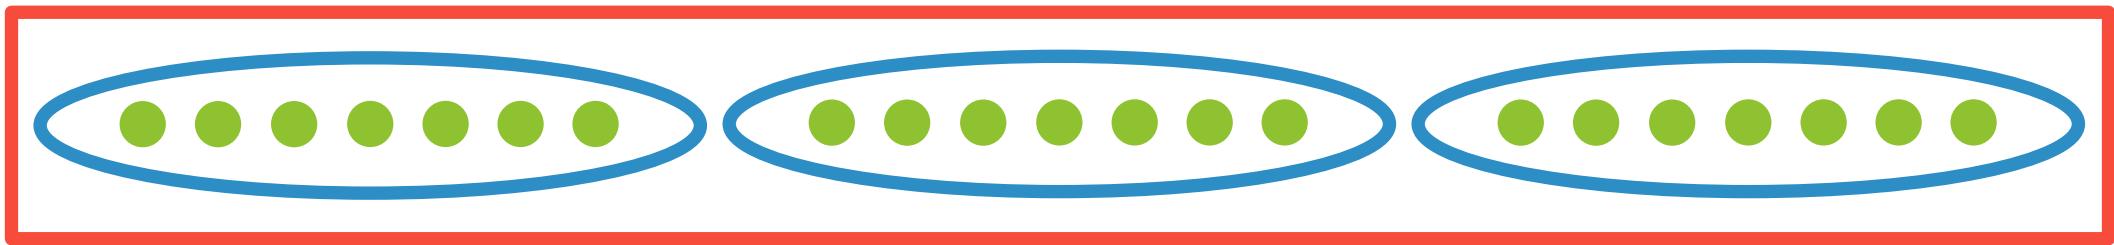
Beautiful Graphs



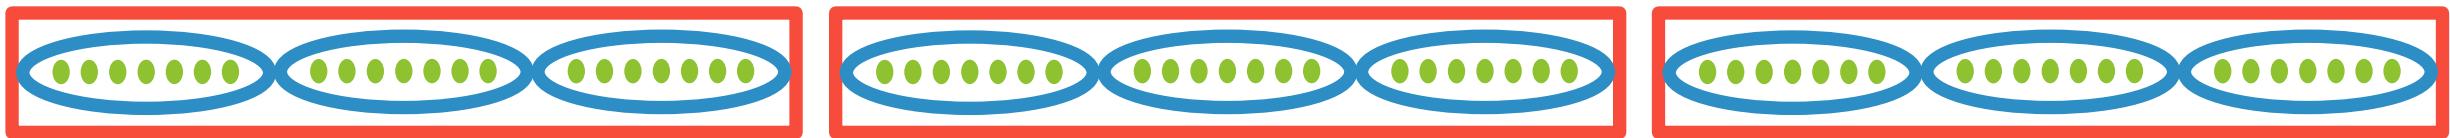
Examples: Cluster Hammer

#EUec08 #sparkbench

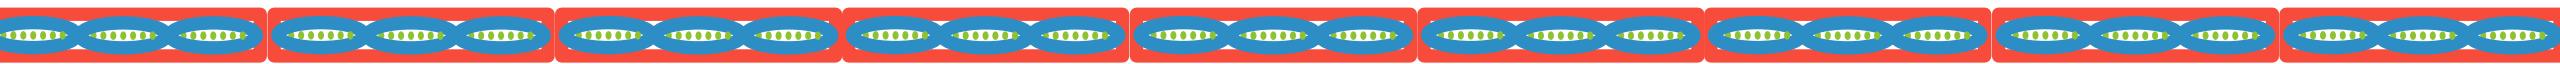
Cluster Hammer



Cluster Hammer



Cluster Hammer

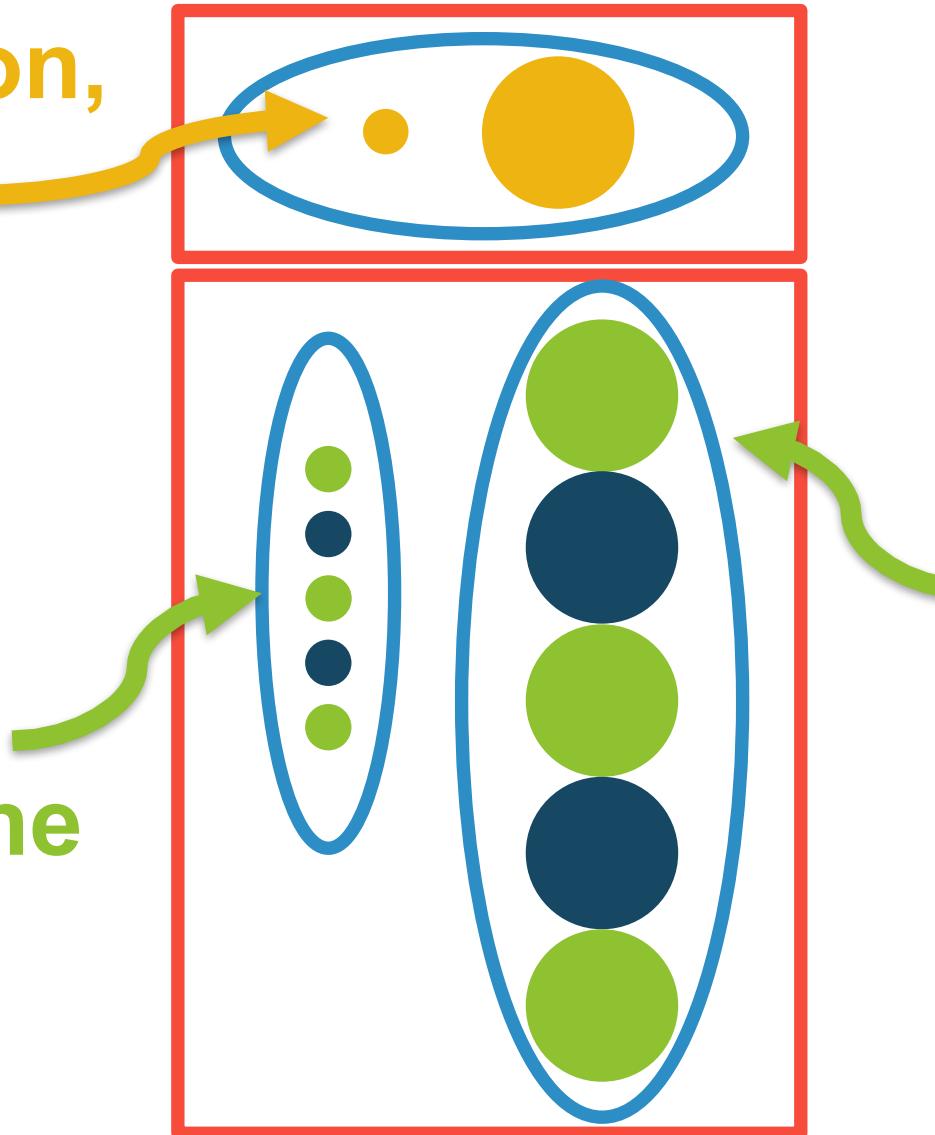


Examples: Multi-User Notebooks

Multi-User Notebook

Data Generation,
sample.csv and
GIANT-SET.csv

Workloads,
working with the
sample.csv



Sleep Workload → 

```
{  
    name = "sleep"  
    distribution = "poisson"  
    mean = 30000  
}
```

Same workloads,
working with
GIANT-SET.csv

Examples of Easy Extensions

- Add more "users" (workload suites)
- Have a workload suite that simulates heavy batch jobs while notebook users attempt to use the same cluster
- Toggle settings for dynamic allocation, etc.
- Benchmark a notebook user on a heavily stressed cluster (cluster hammer!)

Examples: Varying Spark Parameters

Different Versions of Spark

```
spark-bench = {  
    spark-submit-parallel = false  
    spark-submit-config = [ {  
        spark-home = [  
            "/opt/official-spark-build/",  
            "/opt/my-branch-of-spark-with-new-changes/"  
        ]  
        spark-args = {  
            master = "yarn"  
        }  
        workload-suites = [ { // workload suites } ]  
    } ]  
}
```

This will produce 2 spark-submits

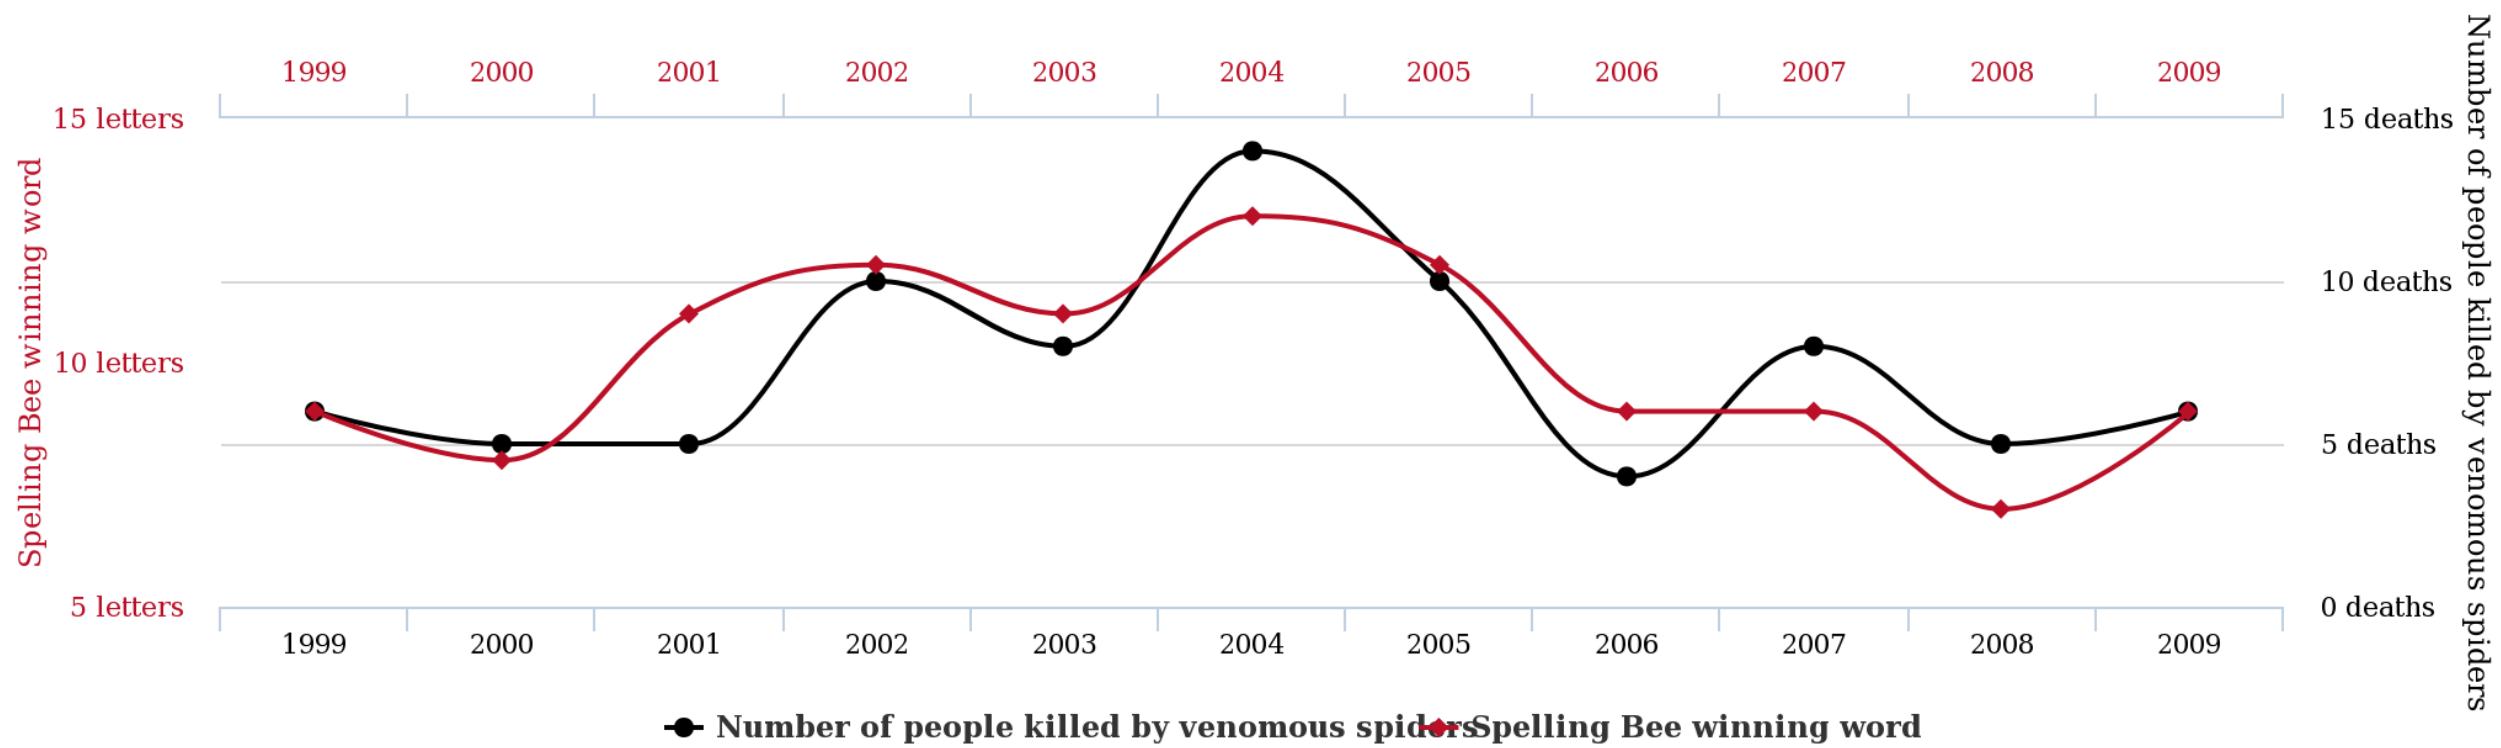
Varying More Than One Param

```
spark-bench = {  
    spark-submit-parallel = false  
    spark-submit-config = [ {  
        spark-home = [  
            "/opt/official-spark-build/",  
            "/opt/my-fork-of-spark-with-new-changes/"  
        ]  
        spark-args = {  
            master = "yarn"  
            executor-mem = ["2G", "4G", "8G", "16G"]  
        }  
        workload-suites = [ { // workload suites } ]  
    } ]  
}
```

This will produce $2 \times 4 = 8$ spark-submits

Beautiful Graphs

Letters in Winning Word of Scripps National Spelling Bee
correlates with
Number of people killed by venomous spiders



Final Thoughts

#EUec08 #sparkbench

Installation

1. Grab the latest release from the [releases page on Github](#)
2. Unpack the tar
3. Set the environment variable for \$SPARK_HOME.
4. Run the examples!

```
./bin/spark-bench.sh examples/minimal-example.conf
```

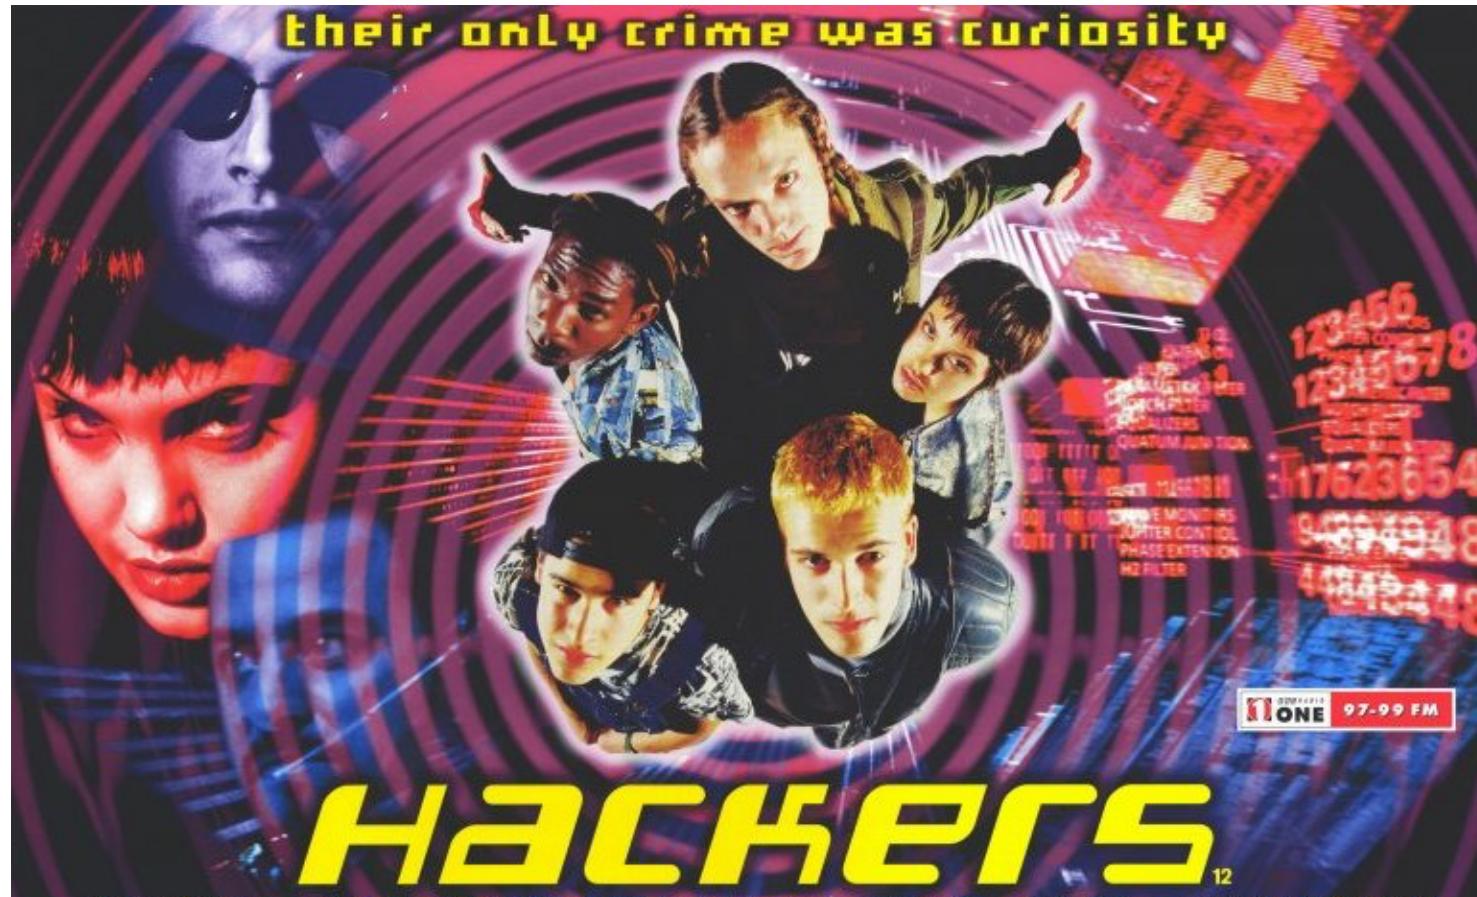
Future Work & Contributions

- More workloads and data generators
- TPC-DS (on top of existing SQL query benchmark)
- Streaming workloads
- Python

Contributions very welcome from all levels!

Check out our ZenHub Board for details

Shout-Out to My Team



*not an actual photo of my team

Matthew Schauer
@showermat

Craig Ingram
@cin

Brad Kaiser
@brad-kaiser

Summary

Spark-Bench is a flexible, configurable framework for
Testing, Benchmarking, Simulating, Comparing,
and more!

<https://sparktc.github.io/spark-bench/>



That's all Folks!

SparkTC/spark-bench

Emily May Curtin

IBM Watson Data Platform
Atlanta, GA



@emilymaycurtin

#EUeco8



@ecurtin

#sparkbench

Project

<https://github.com/SparkTC/spark-bench>

Documentation

<https://sparktc.github.io/spark-bench/>