



Real-time Detection Of Anomalies In The Database Infrastructure Using Apache Spark

Prasanth Kothuri, CERN

Daniel Lanza Garcia, CERN

#EUde13

Outline

- About CERN
- Apache Spark at CERN
- Data Lake Challenges
- Detecting Anomalous Connections
- Real-time Intelligent Monitoring

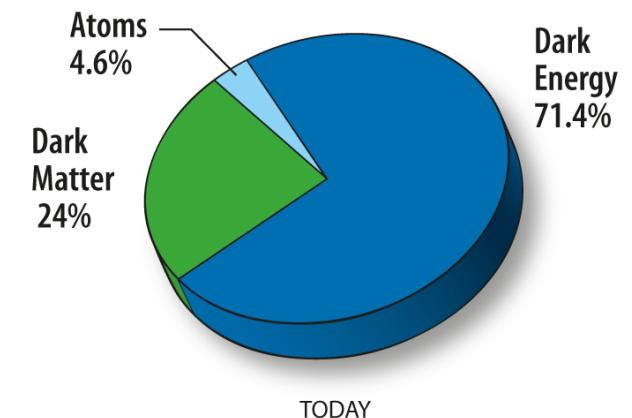
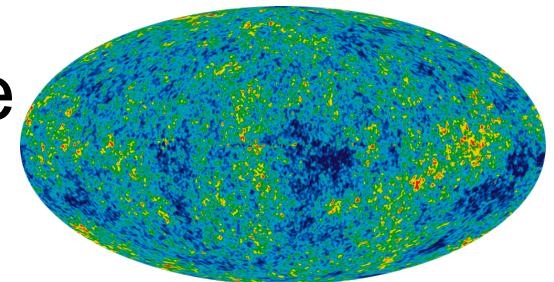
About CERN

- CERN - European Laboratory for Particle Physics
- Founded in 1954 by 12 Countries for fundamental physics research in the post-war Europe
- Today 21 member states + world-wide collaborations
 - About ~1000 MCHF yearly budget
 - 2'300 CERN personnel
 - 10'000 users from 110 countries



Fundamental Research

- What is 95% of the Universe made of?
- Why do particles have mass?
- Why is there no antimatter left in the Universe?
- What was the Universe like just after the “Big Bang”



The Large Hadron Collider (LHC)



Largest machine in the world

27km, 6000+ superconducting magnets

Fastest racetrack on Earth

Protons circulate 11245 times/s (99.9999991% the speed of light)

Emptiest place in the solar system

High vacuum inside the magnets

Hottest spot in the galaxy

During Lead ion collisions create temperatures 100 000x hotter than the heart of the sun;

Apache Spark @ CERN

- Spark is a popular compute engine for Analytics + ETL
 - Deployed on four production Hadoop/YARN clusters
 - Aggregated capacity (2017): ~1500 physical cores, 11 PB
 - Adoption is growing. Key projects involving Spark:
 - Analytics for **accelerator** controls and logging
 - Monitoring **IT Infrastructure**, uses Spark + Spark streaming
 - **Analytics** on aggregated logs
 - Explorations on the use of Spark for **high energy physics**

Detection Of Anomalies In Database Infrastructure

CERN Database Infrastructure

- Over 100 Oracle databases, most of them RAC
 - Running Oracle 11.2.0.4 and **12.1.0.2**
 - 950 TB of data files for production DBs in total (without counting replicas), NAS as storage
 - Oracle In-Memory in production since July 2015
- Examples of critical production DBs:
 - Quench Protection System: 150'000 changes/s
 - LHC logging database: 492 TB, growing 180 TB/year
- But also DB as a Service (single instances)
 - 310 MySQL, 70 PostgreSQL, 9 Oracle, 5 InfluxDB

ORACLE



InfluxDB

MySQL

PostgreSQL

Data Lake

- Objective: Build central repository for database audit, performance metrics and logs with the goal of real time analytics and offline analytics
- Use cases
 - Detection of Anomalies
 - Troubleshooting
 - Capacity Planning

Diverse Data

- Tables
 - Audit
 - Performance Metrics
 - Alert
- Log files
 - Listener

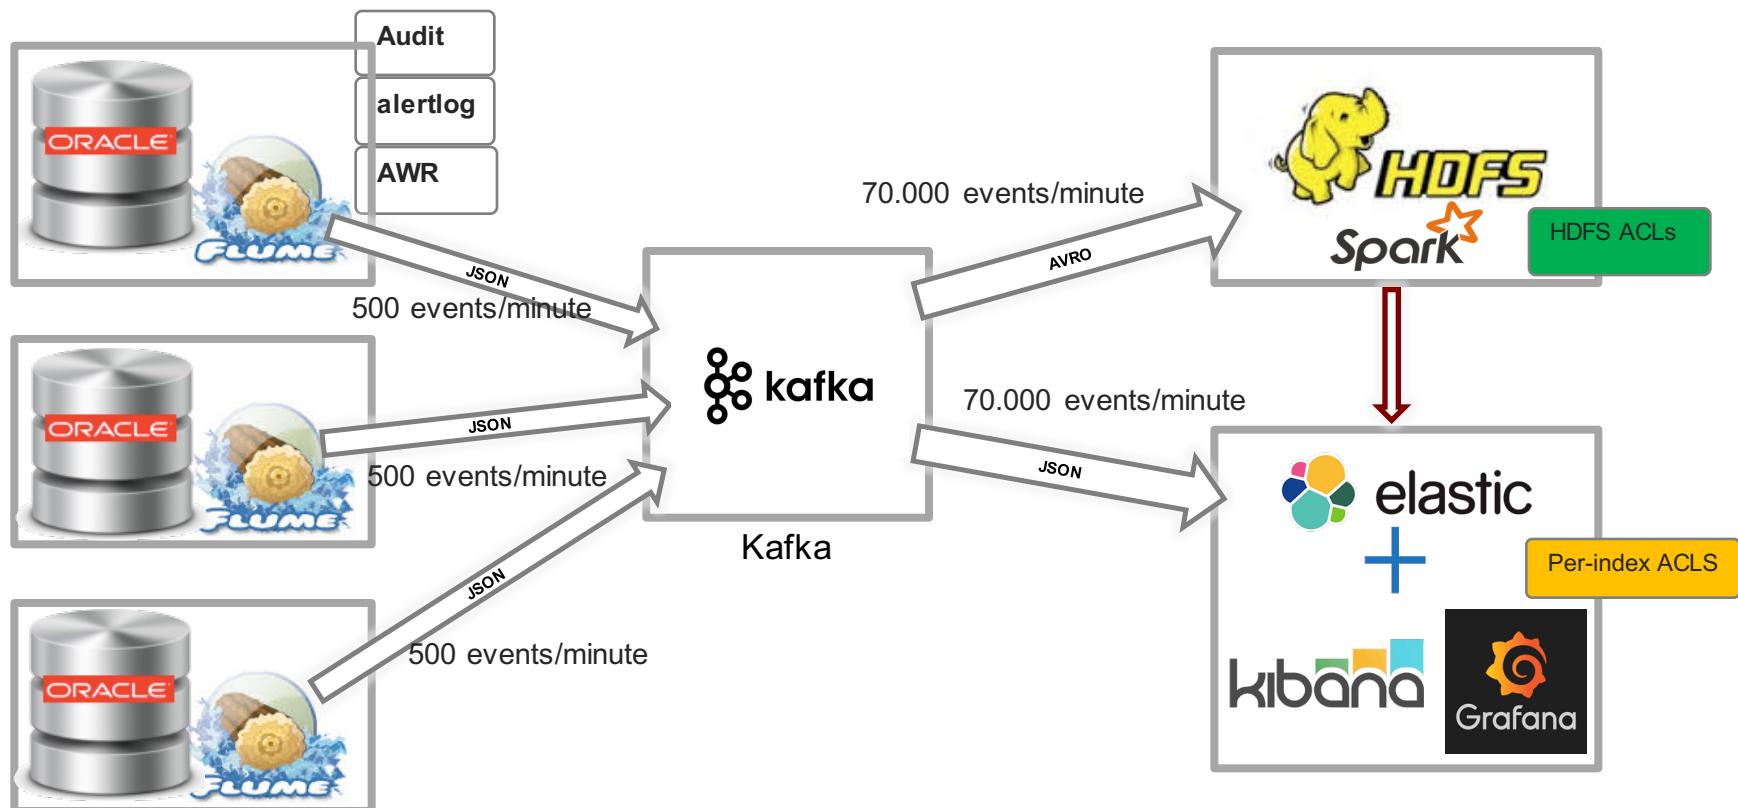
```
oracle_sid: CMSINTR2 database_type: oracle source_type: alert hostname: cmsrac44 flume_agent_version: 0.1.4-5.el6 database_version: 11.2.0.4.0 ADDR: 00007F80FA4511E0 INDX: 31,991 INST_ID: 2 ORIGINATING_TIMESTAMP: November 9th 2016, 09:50:17.000 NORMALIZED_TIMESTAMP: - ORGANIZATION_ID: oracle COMPONENT_ID: rdbms HOST_ID: cmsrac44.cern.ch HOST_ADDRESS: 10.176.84.61 MESSAGE_TYPE: 1 MESSAGE_LEVEL: 16 MESSAGE_ID: - MESSAGE_GROUP: - CLIENT_ID: - MODULE_ID: - PROCESS_ID: 15031 THREAD_ID: - USER_ID: - INSTANCE_ID: - DETAILED_LOCATION: -
```

```
oracle_sid: C2PPSDB2 database_type: oracle producer: oracle source_type: metric hostname: itrac1329.cern.ch flume_agent_version: 0.1.6-7.el6 type: dbconnection database_version: 11.2.0.4.0 DB_NAME: C2PPSDB DB_UNIQUE_NAME: C2PPSDB_RAC13 INSTANCE_NAME: C2PPSDB2 BEGIN_TIME: June 29th 2017, 11:51:11.000 END_TIME: June 29th 2017, 11:52:11.000 METRIC_NAME: Enqueue Deadlocks Per Txn VALUE: 0B _id: AVzzRAX-Nj0vFzICP54W _type: log _index: itdb_db-metric-2017-06-29 _score: -
```

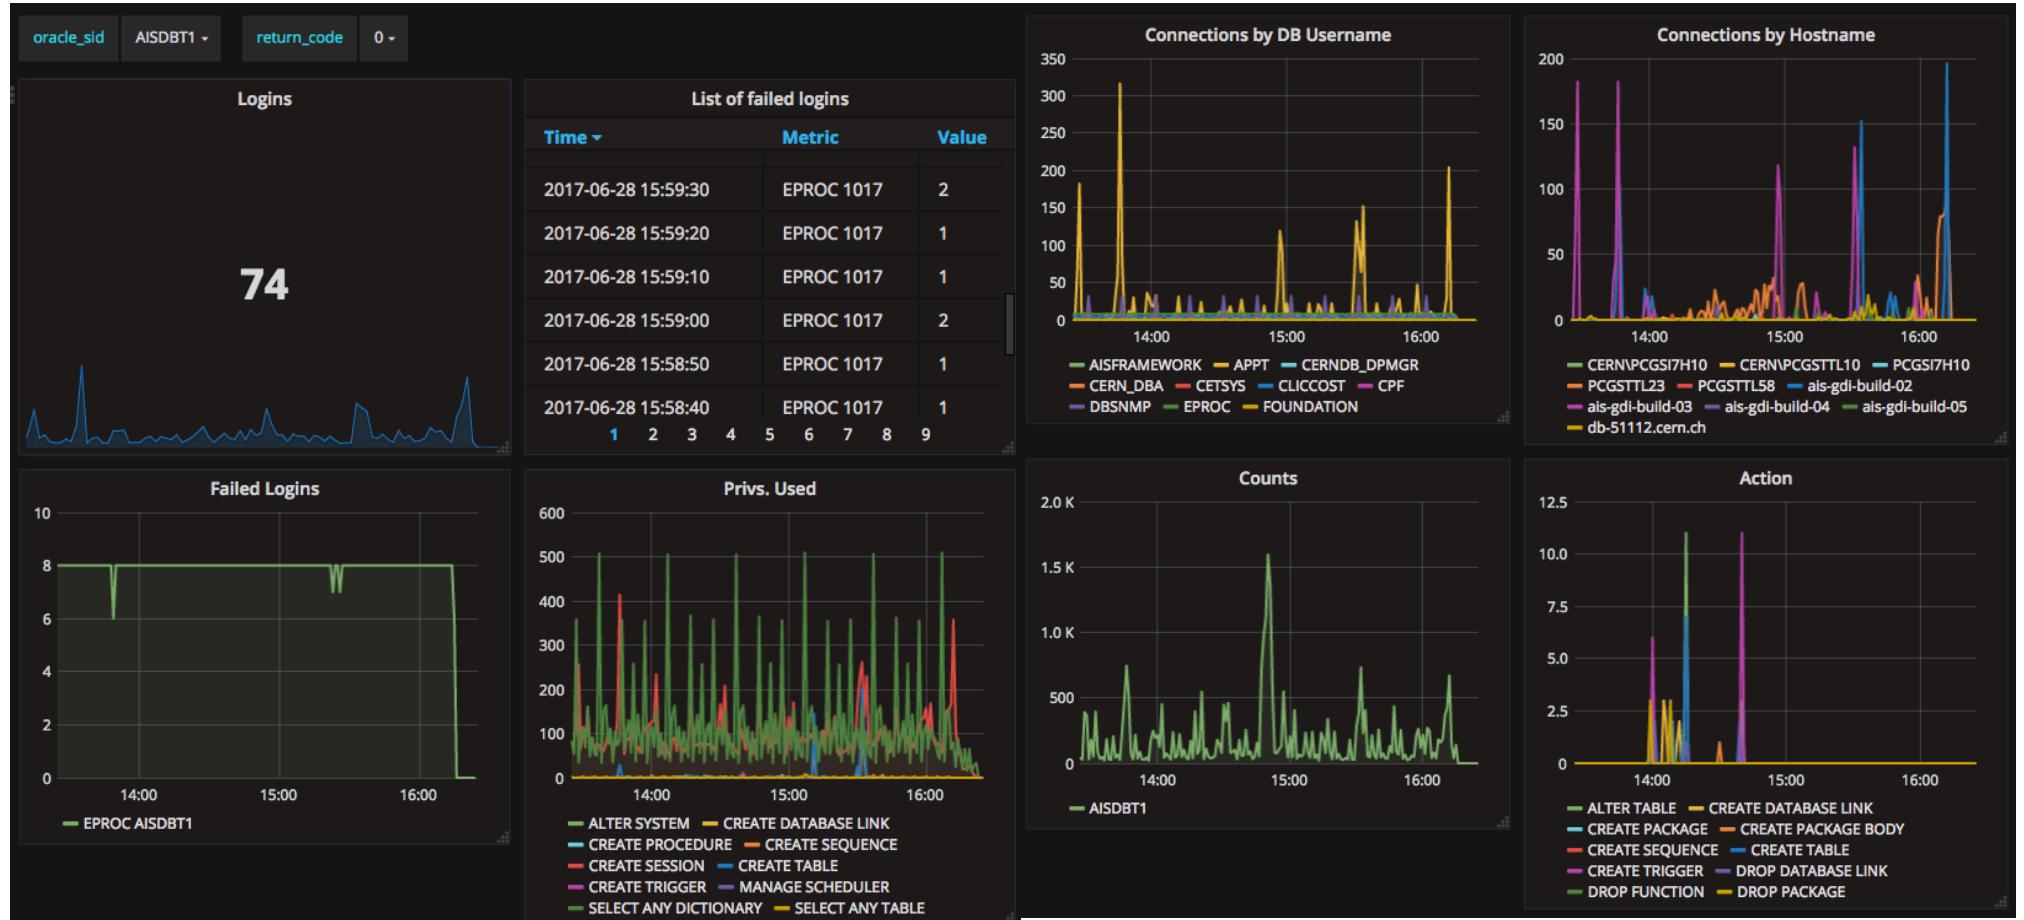
```
29-JUL-2016 15:17:34 * (CONNECT_DATA=(SID=DESFONUD)(CID=(PROGRAM=oracle)(HOST=itrac50035.29-JUL-2016 15:17:38 * service_update * WCERND * 0
Fri Jul 29 15:18:45 2016
29-JUL-2016 15:18:45 * (CONNECT_DATA=(SID=PAYT)(CID=(PROGRAM=RTSDGN@db-50016)(HOST=db-500
En: 1.1 20 15.10.51 2016
```

t client_host	db-51220.cern.ch
t client_ip	137.138.161.92
# client_port	33,450
t client_program	perl
t client_protocol	tcp
t client_user	sysctl

Data pipeline Architecture



Aggregations and Dashboards



Real-time analysis

- Attempt to detect anomalous connections using Machine Learning techniques
- First approach to use kNN classifier to classify connection as *normal* or *anomalous*
- Anomaly detection suffer from high false positive rate, the challenge is to choose features that best characterize user connections pattern

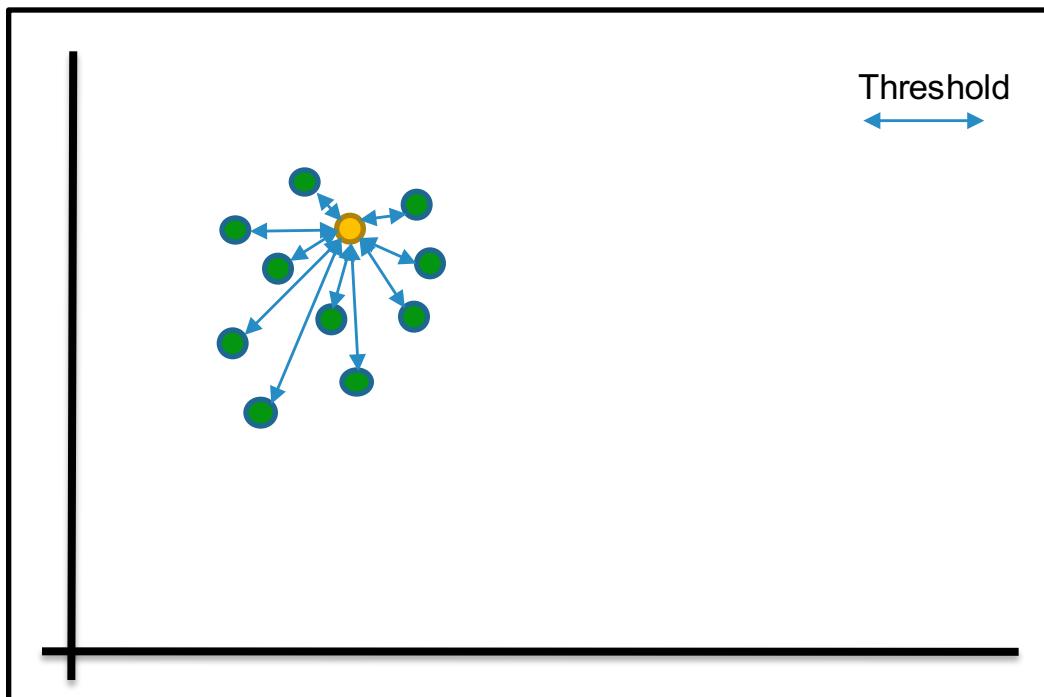
Feature Selection

- Data Preparation
- Feature Extraction
 - Client_host
 - Client_user
 - Client_program
 - Service_name
 - Hour_of_the_day
 - Day_of_the_week

```
{  
  "oracle_sid": "XXXX_YYYY",  
  "listener_name": "listener_scan1",  
  "database_type": "oracle",  
  "producer": "oracle",  
  "source_type": "listener",  
  "hostname": "XXXXXX.cern.ch",  
  "flume_agent_version": "0.1.6-7.el6",  
  "type": "dbconnection",  
  "event_timestamp": "2017-09-27T04:45:27+0200",  
  "CONNECT_DATA_SERVER": "DEDICATED",  
  "CONNECT_DATA_SERVICE_NAME": "XXXr.cern.ch",  
  "client_program": "python",  
  "client_host": "pxxxvj2.cern.ch",  
  "client_user": "merge",  
  "client_protocol": "tcp",  
  "client_ip": "137.100.100.167",  
  "client_port": 38432,  
  "type": "establish",  
  "service_name": "XXXX.cern.ch",  
  "return_code": 0  
}
```

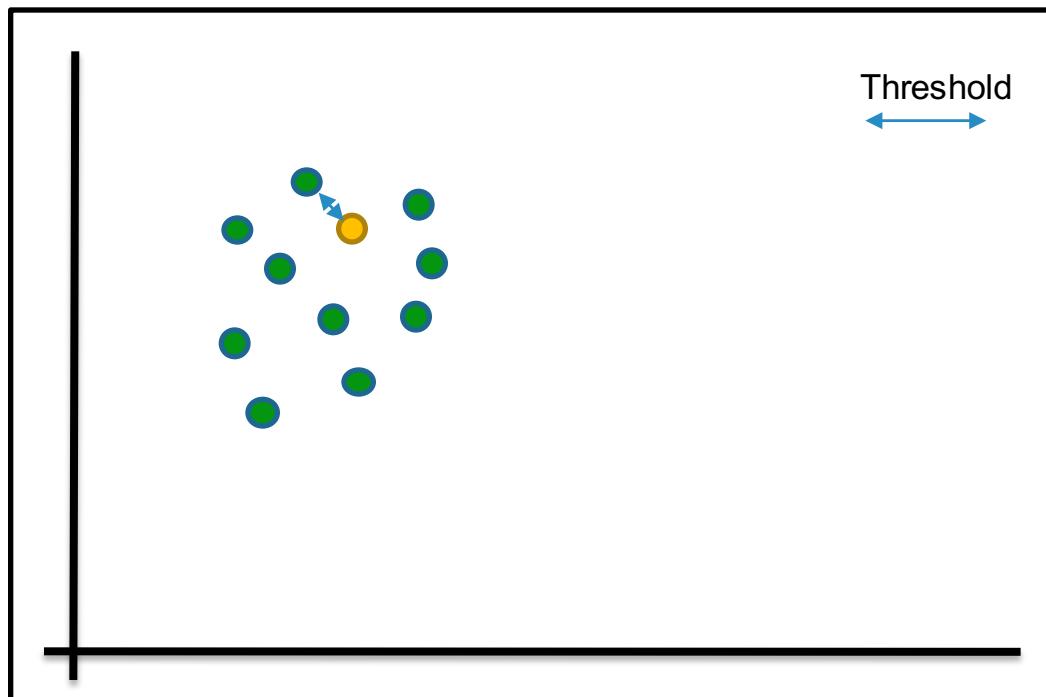
Detecting anomalous connections job

- Machine learning: kNN



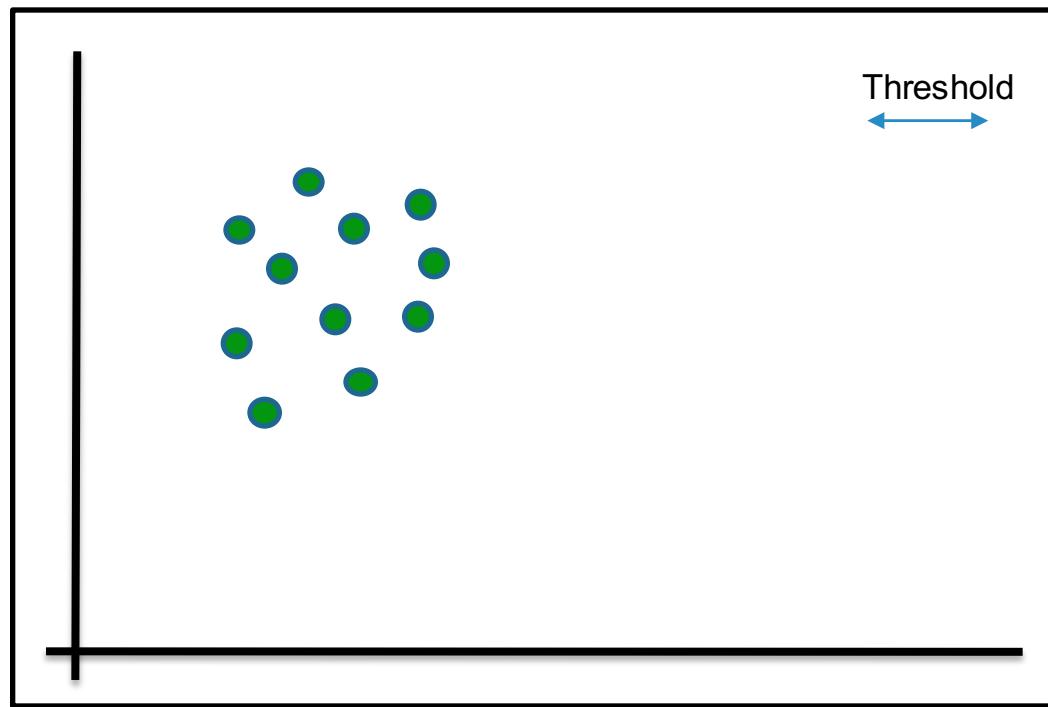
Detecting anomalous connections job

- Machine learning: kNN



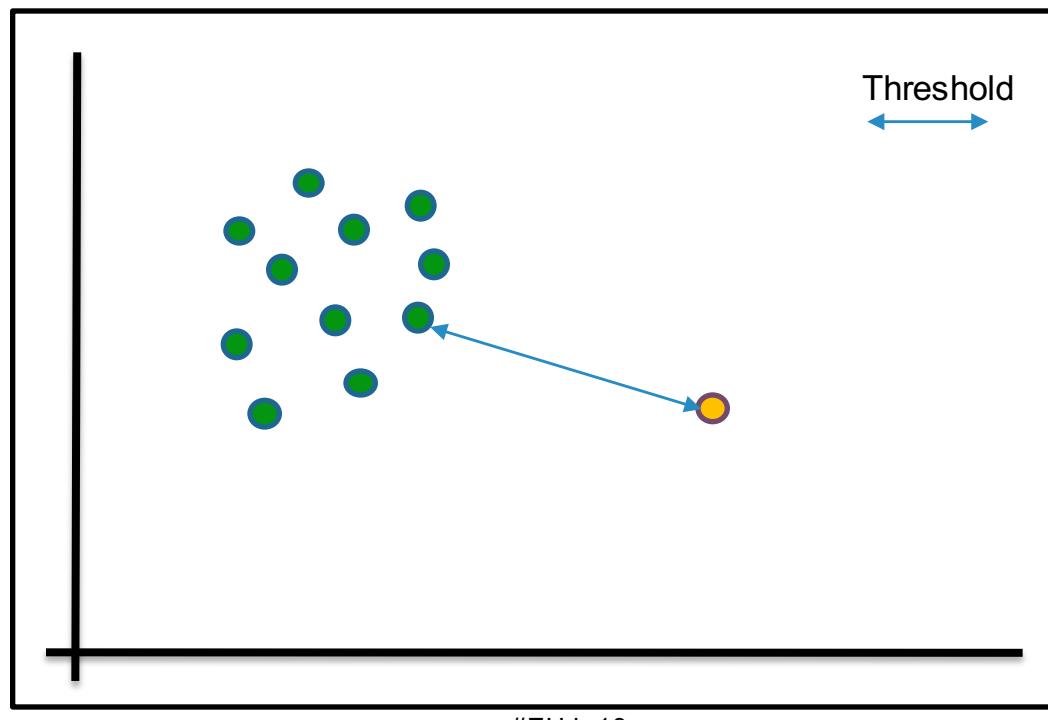
Detecting anomalous connections job

- Machine learning: kNN



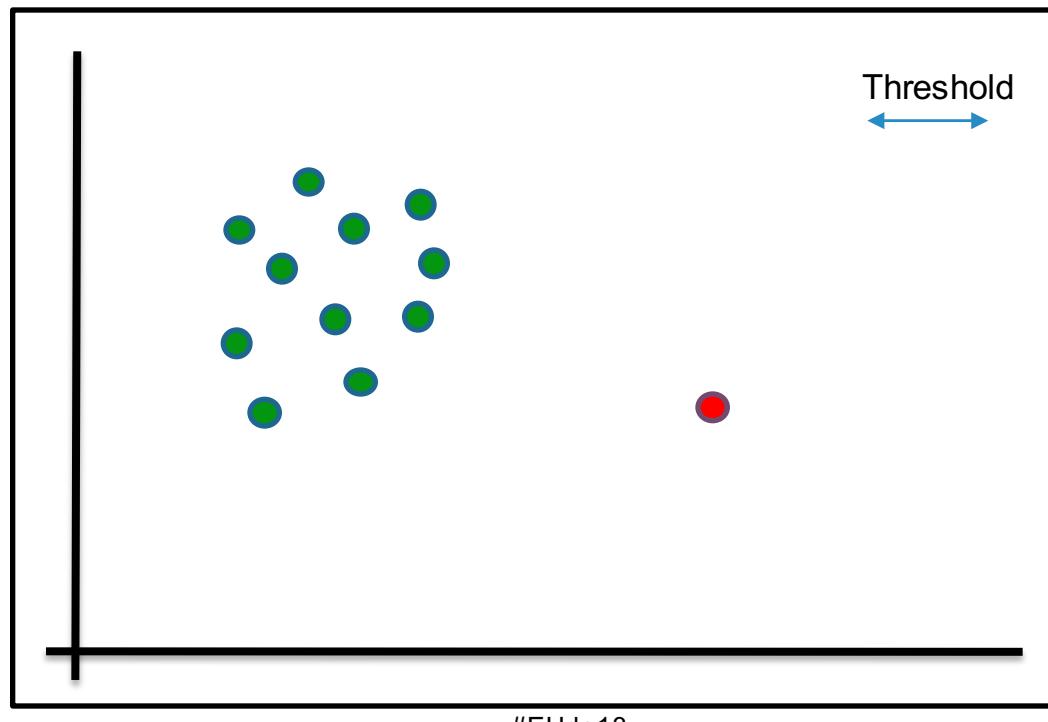
Detecting anomalous connections job

- Machine learning: kNN



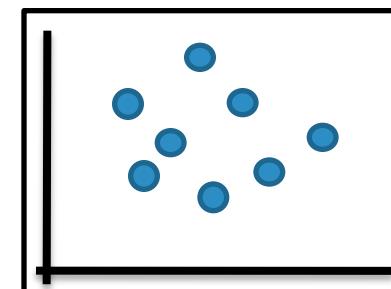
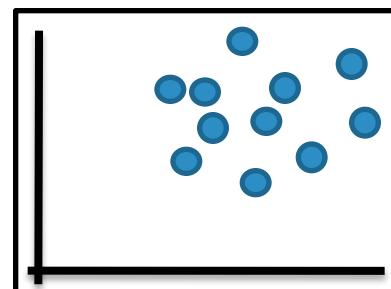
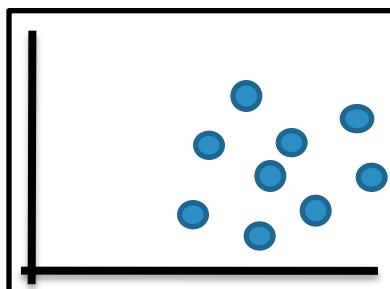
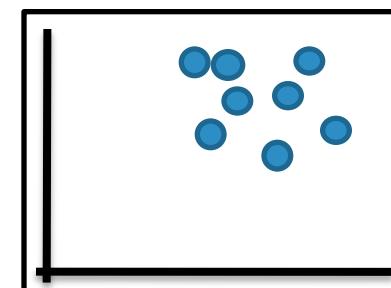
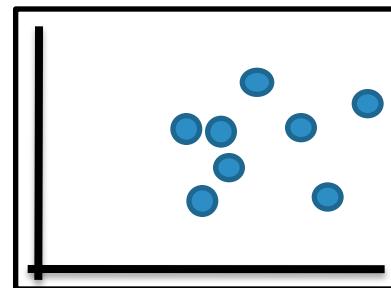
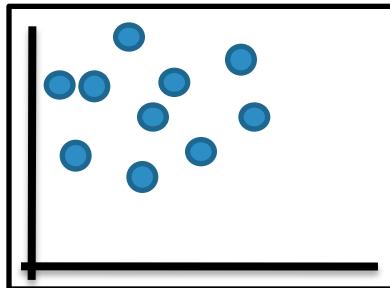
Detecting anomalous connections job

- Machine learning: kNN



Detecting anomalous connections job

- Machine learning: kNN
 - Making it scalable, each user should behave similarly



Detecting anomalous connections job

- Distance Function:

```
for (Connection trustedConnection : connections) {  
    float distance = connection.distance(dataMineProperties, trustedConnection);  
  
    if(distance < minimun_distance)  
        minimun_distance = distance;  
}
```

- Spark Streaming Job processing – 10 million / day
- 6 executors, each with 4 cores and 8GB of memory

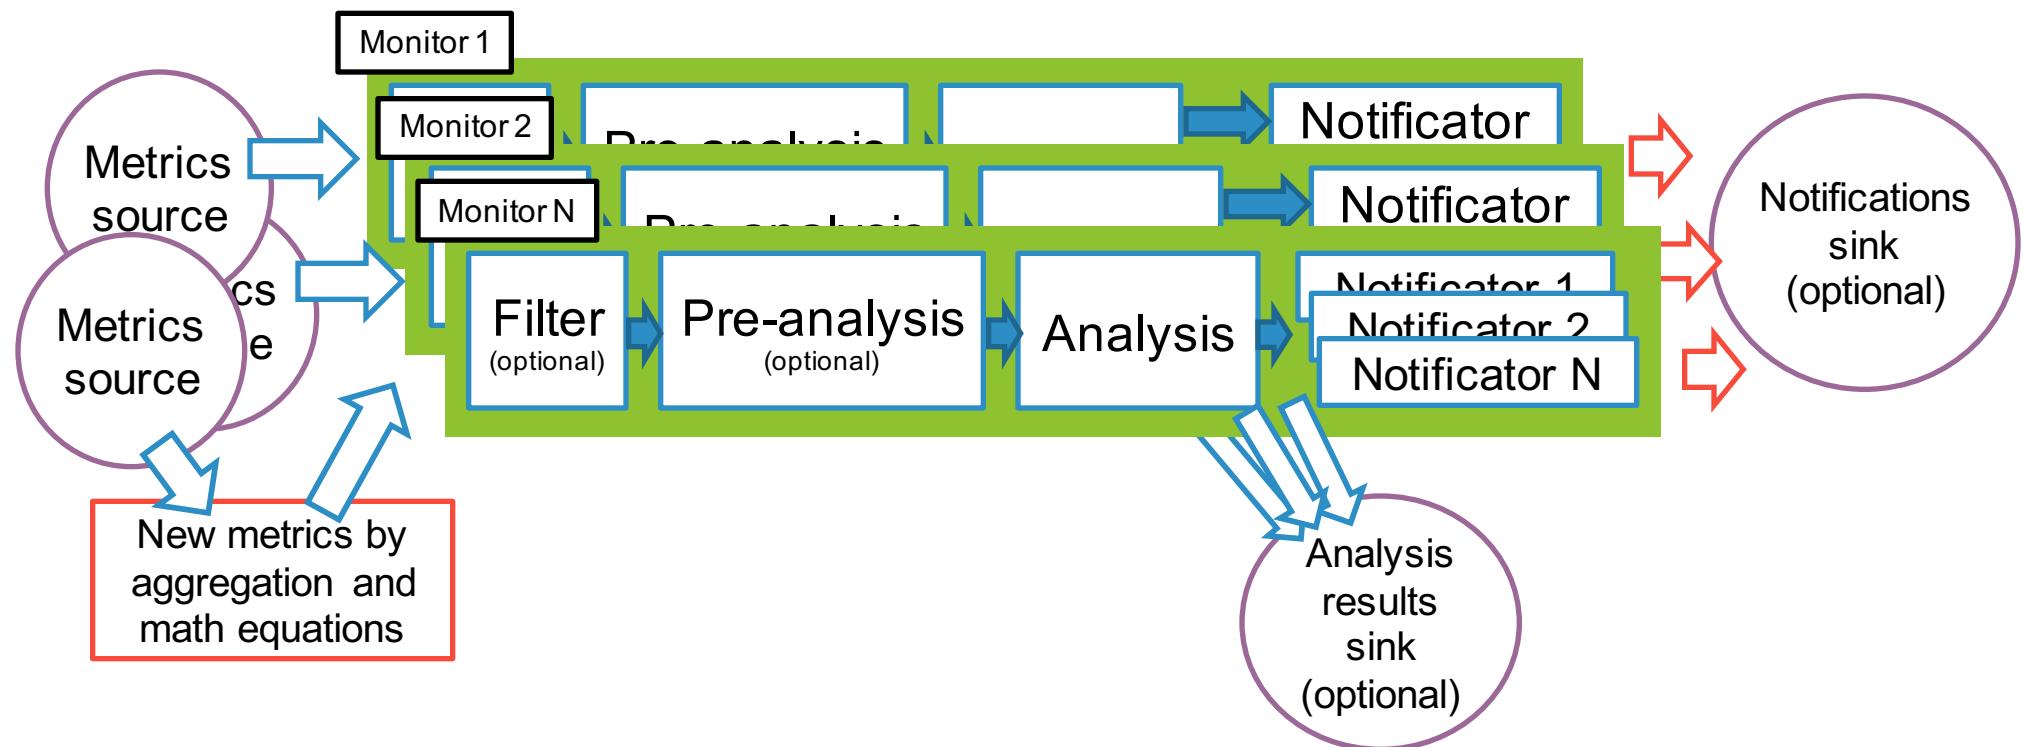
Metrics monitor job

- A general purpose metrics monitor

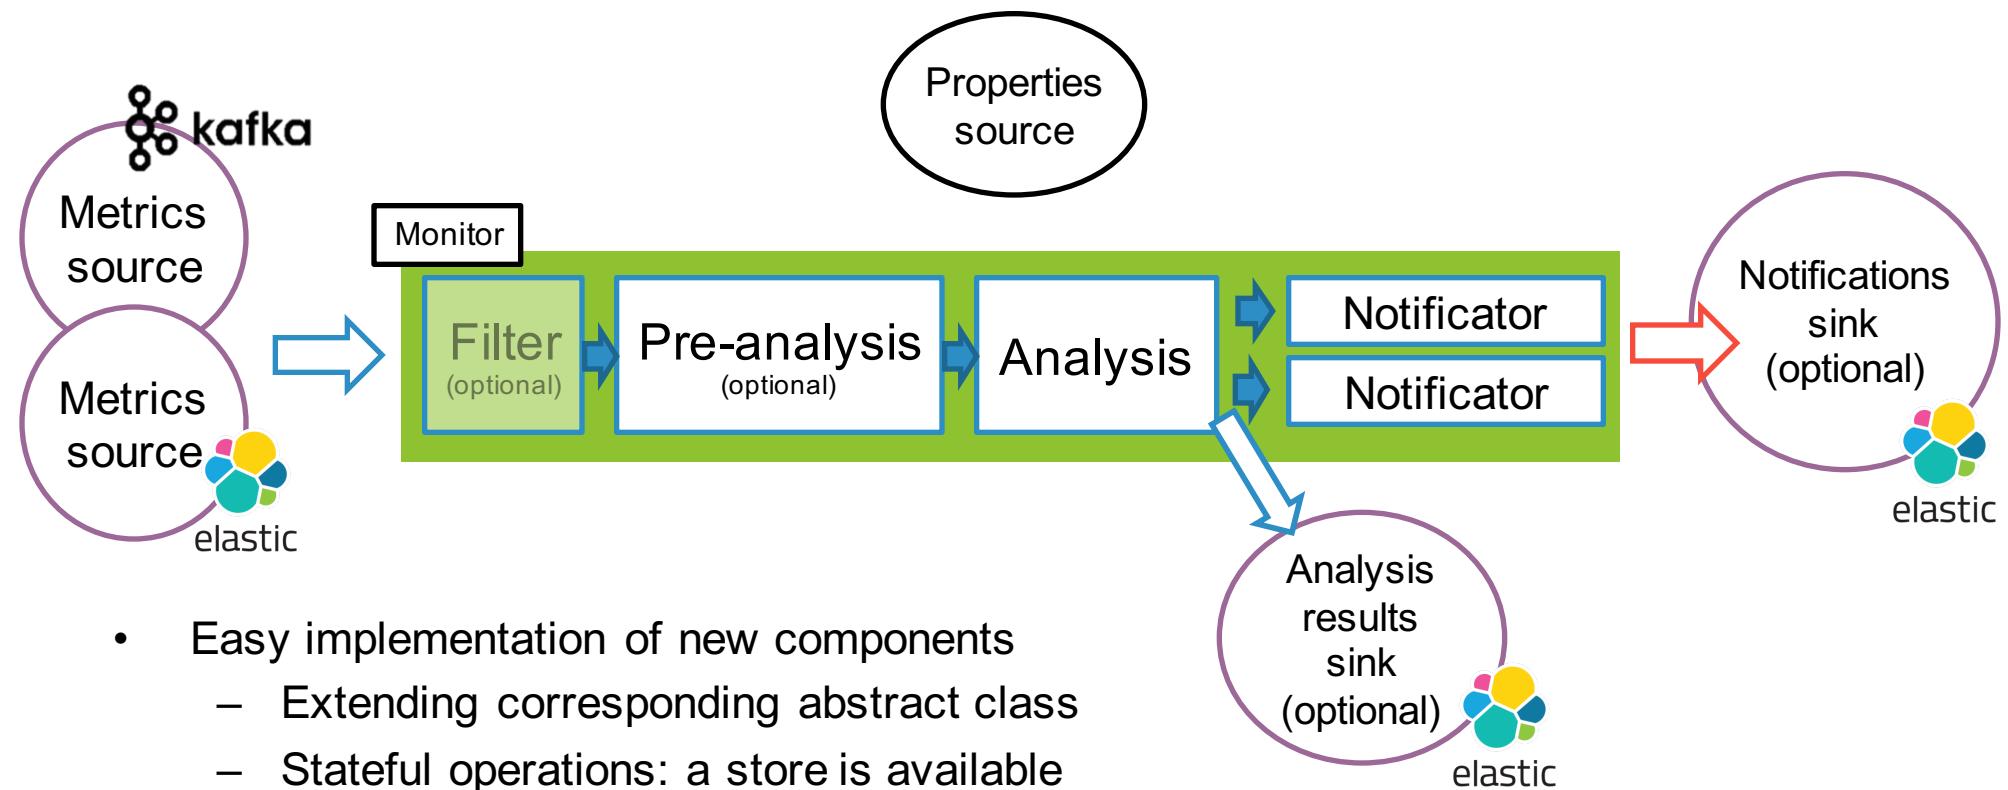


Metrics monitor job: data flow

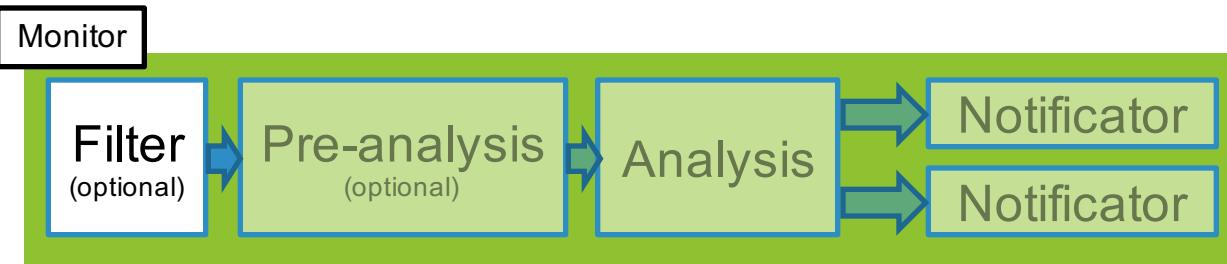
Spark
Streaming



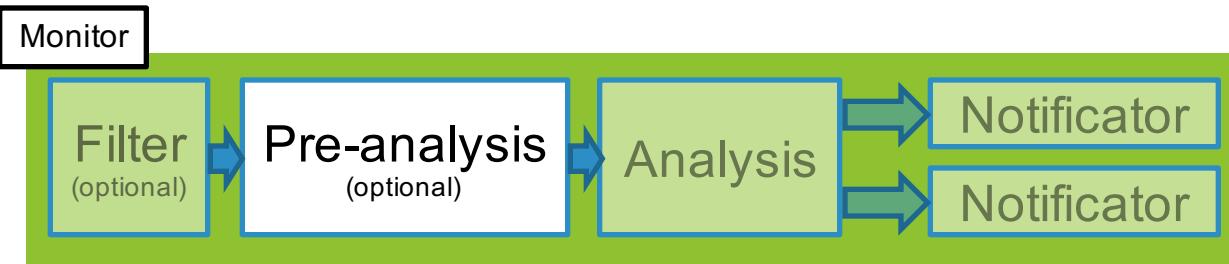
Metrics monitor job: modularity



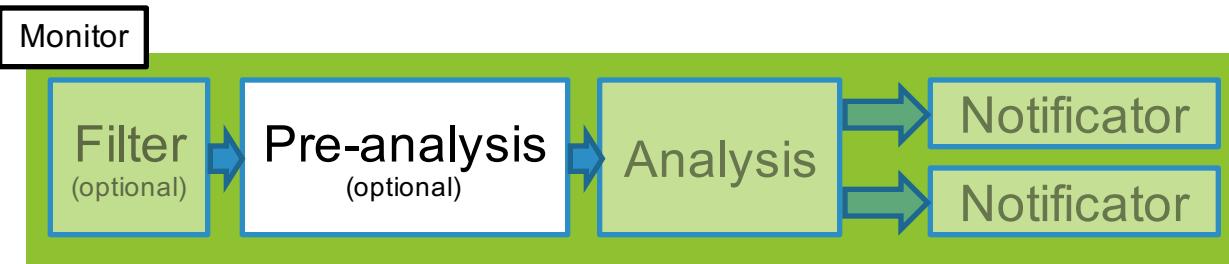
- Easy implementation of new components
 - Extending corresponding abstract class
 - Stateful operations: a store is available



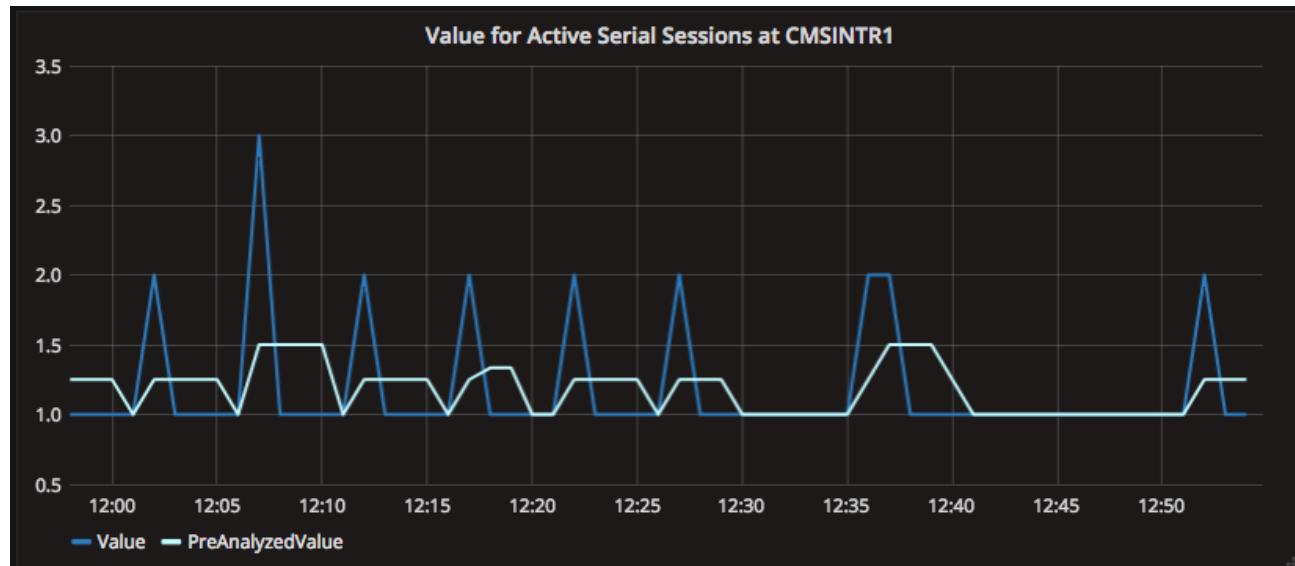
- Which metrics should process this monitor?
 - CPU Usage of all databases, Cache Hit % of MySQL databases, all machines in PROD, ...
- Filter metrics by attributes
- Accepts exact value or regular expressions
- If not applied, all metrics are processed

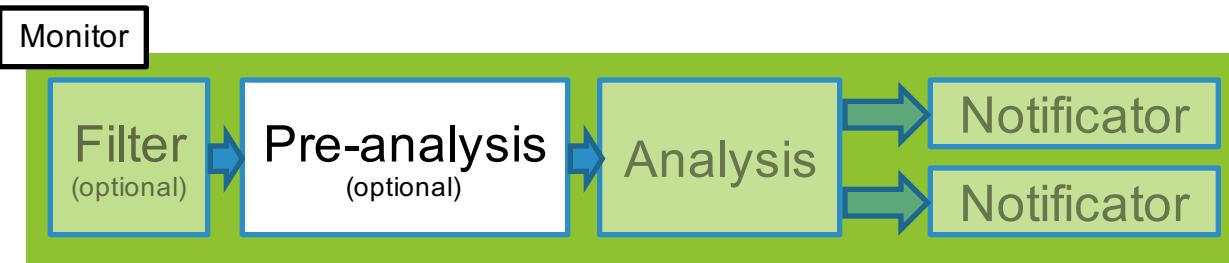


- Message metric value
 - Average, difference, filter, ...
- Internal store for stateful pre-analysis
- If applied, result will be the analyzed value

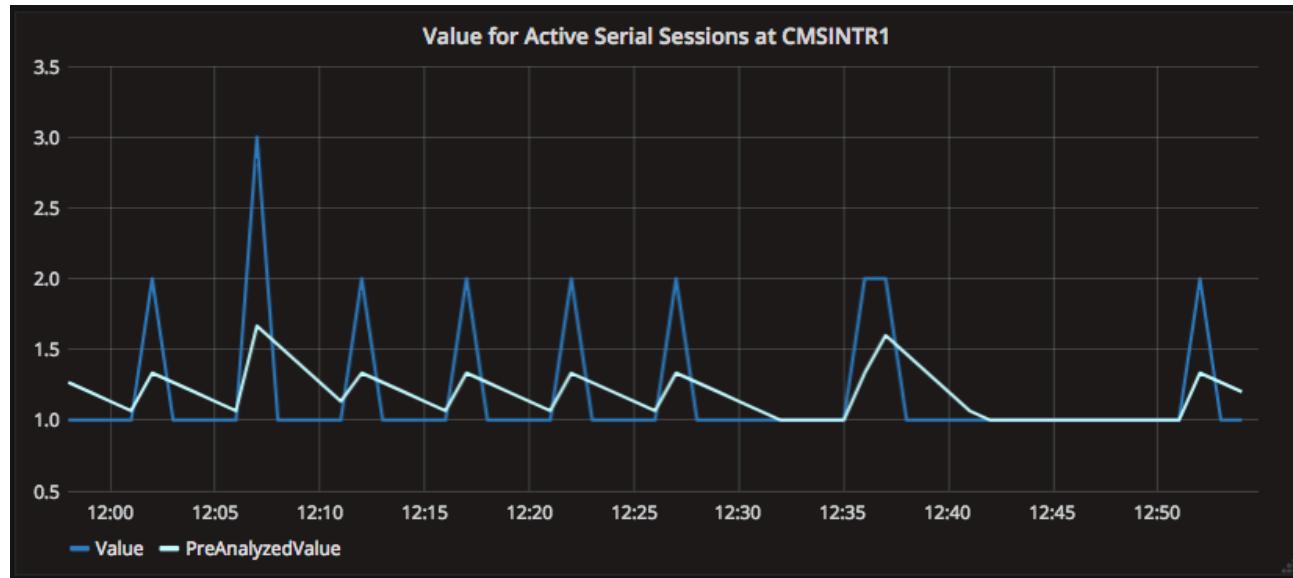


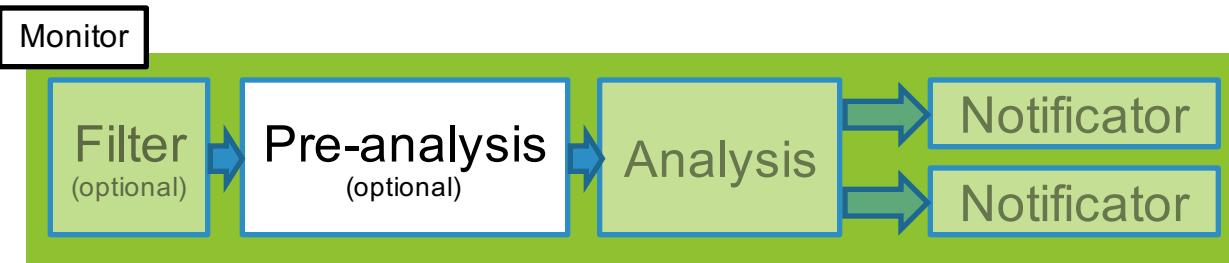
- Type: average value of a period



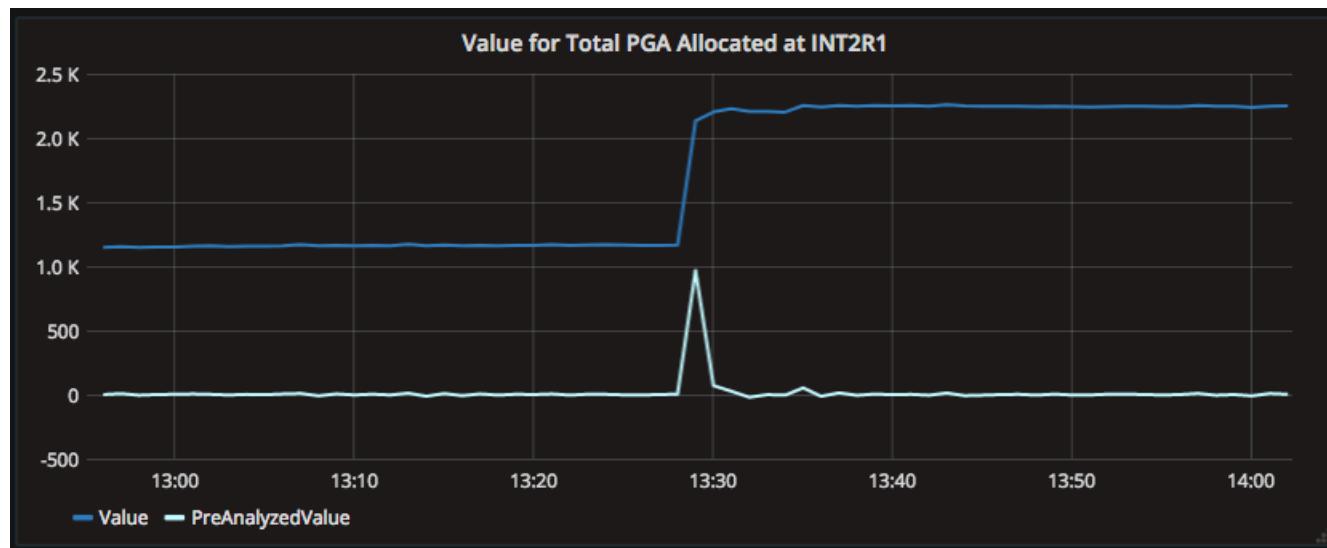


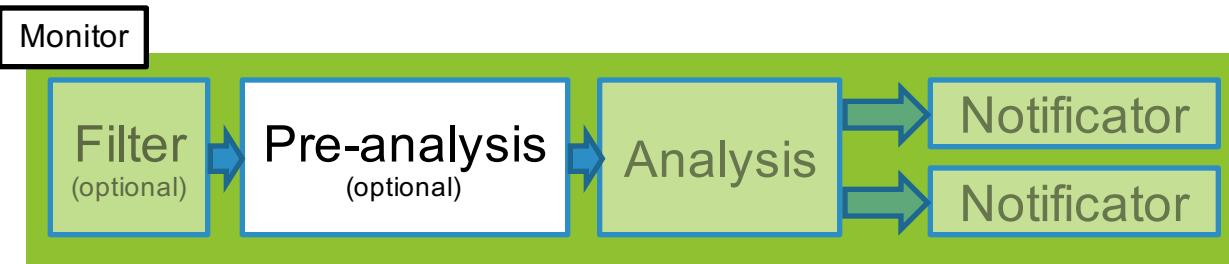
- Type: weighted average value of a period



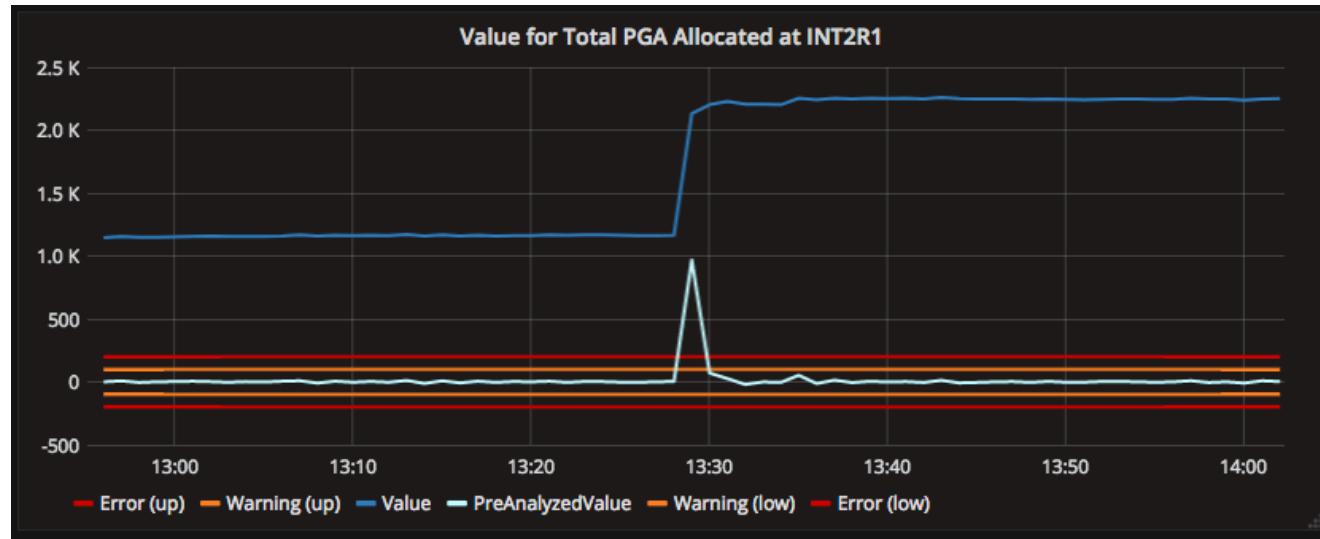


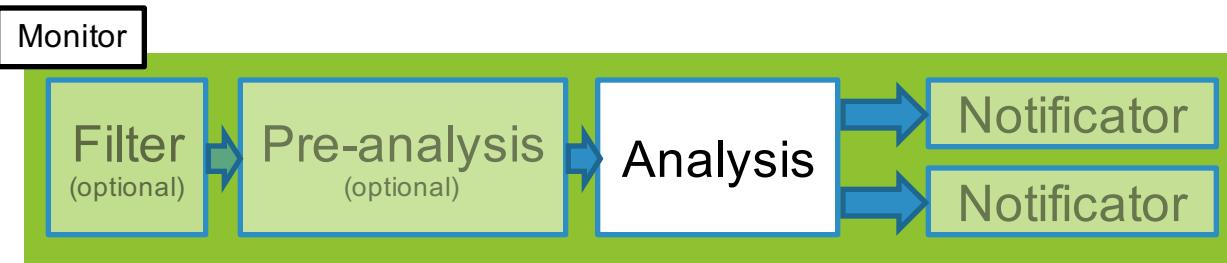
- Type: difference with previous value



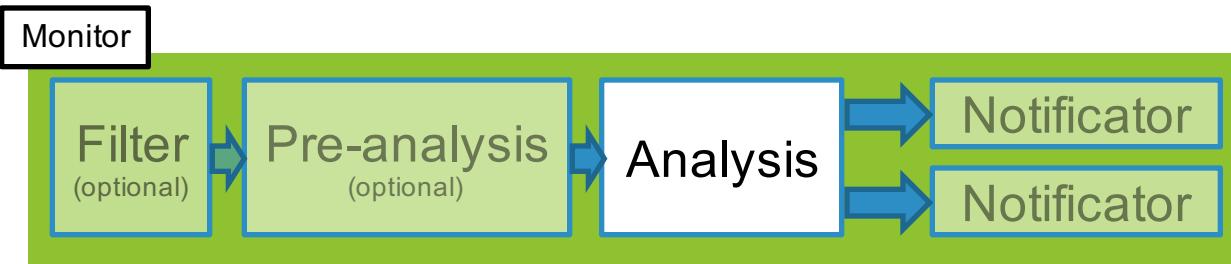


- Type: difference with previous value

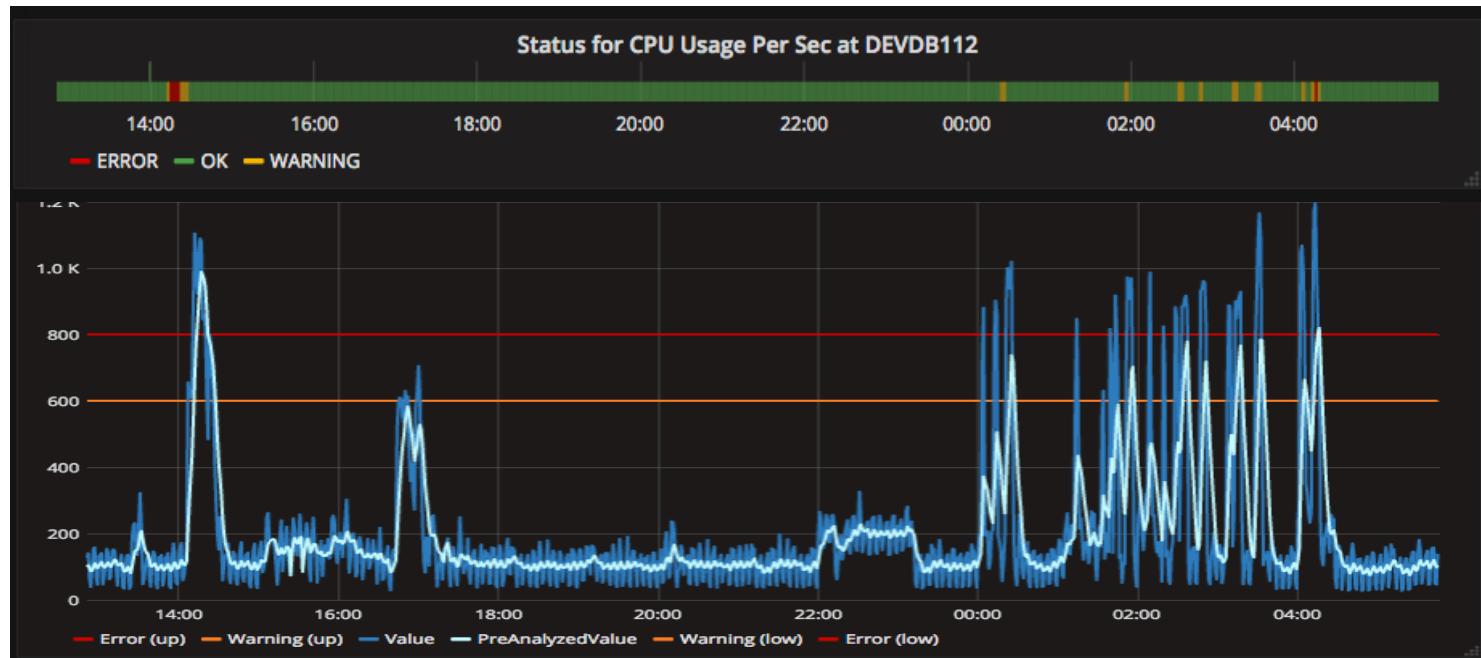




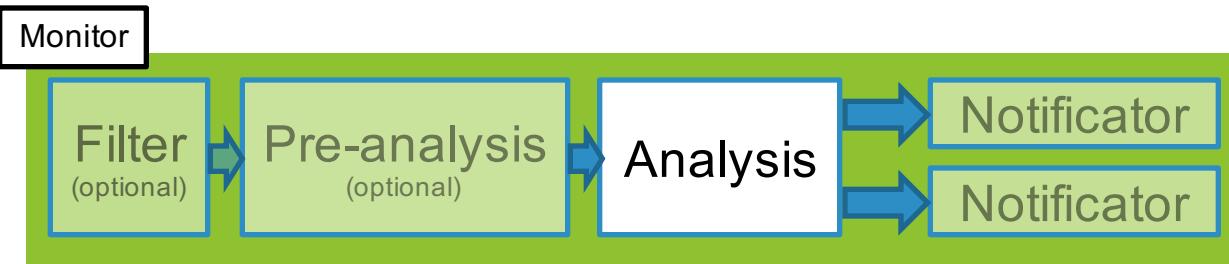
- Is the metric behaving as it should?
- Determine the status of each incoming metric
 - OK, WARNING, ERROR, EXCEPTION
- Results can be sent to external storage
- Internal store for stateful analysis



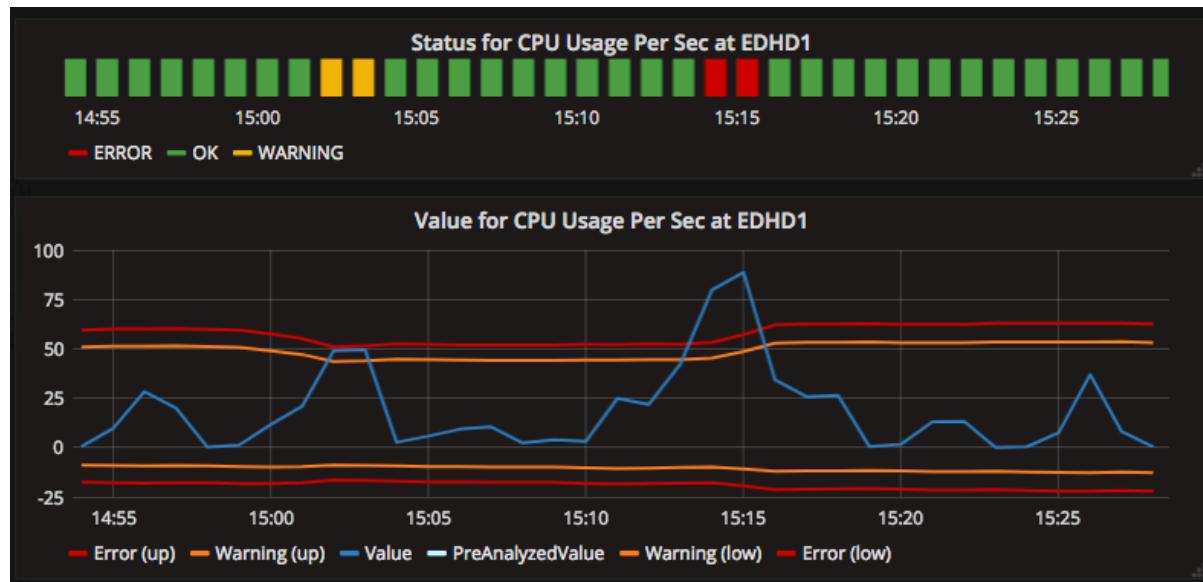
- Type: thresholds fixed manually

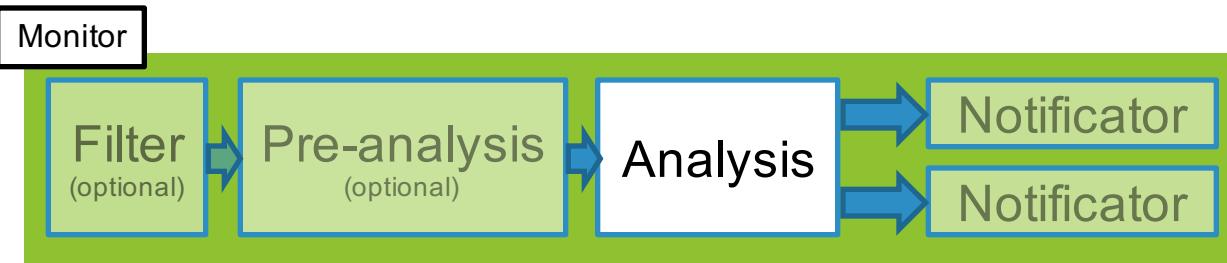


#EUde13

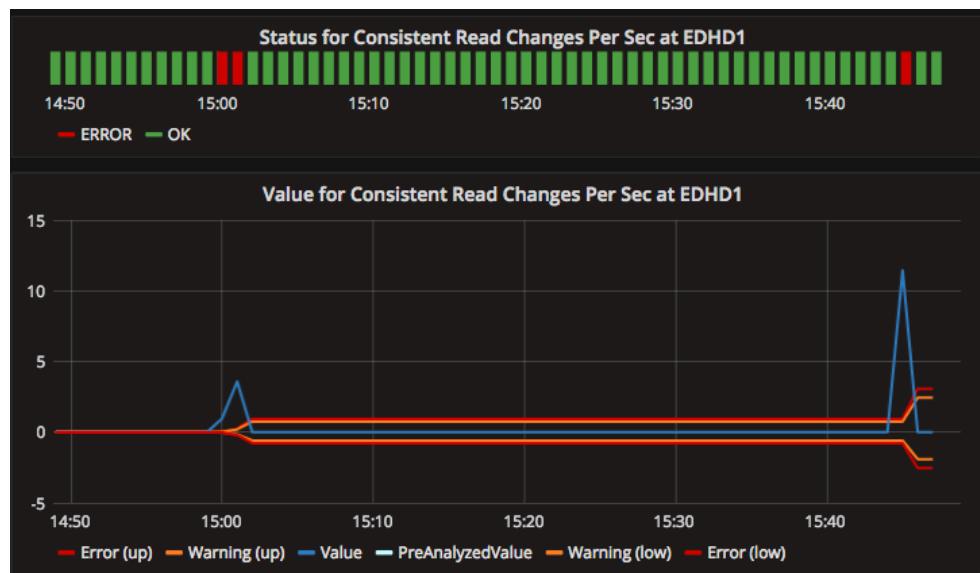


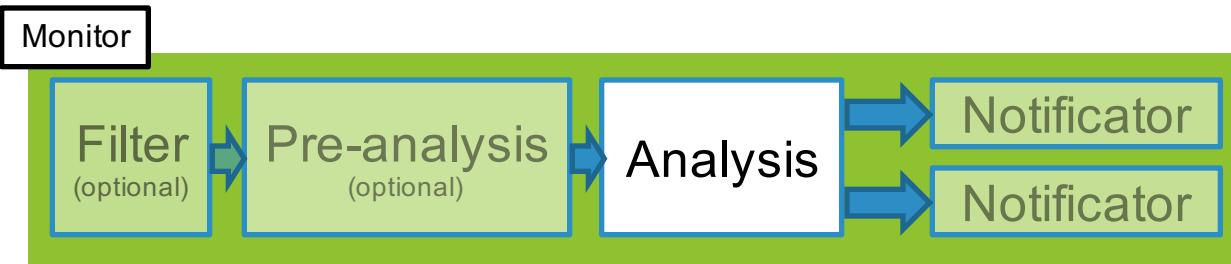
- Type: learning from recent behaviour (average and variance)
 - Errors when metric does not behave as it used to



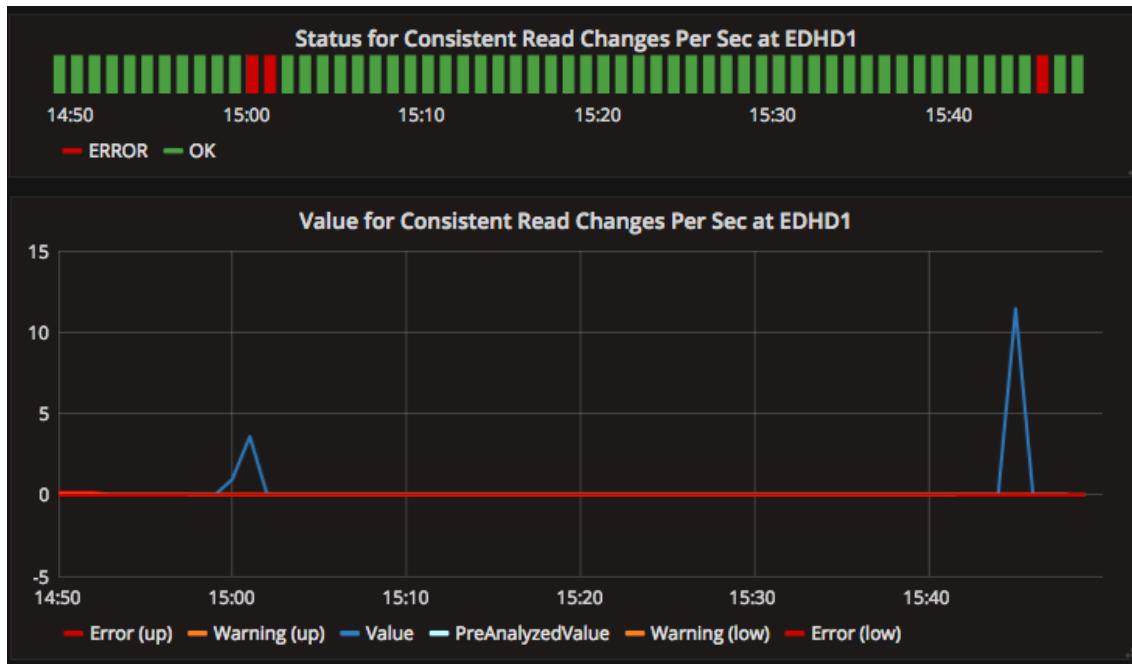


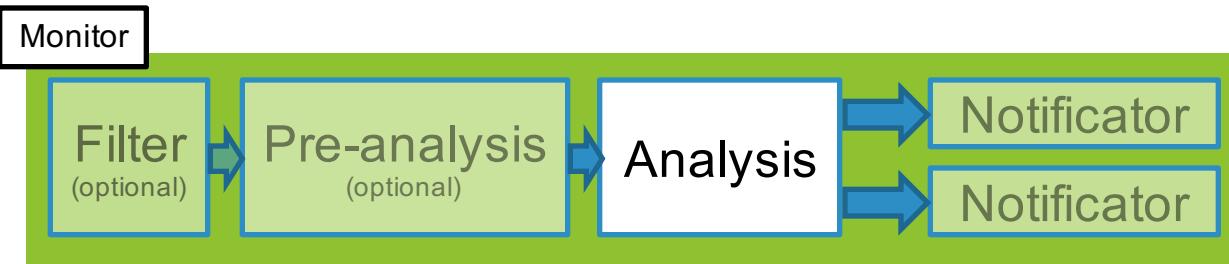
- Type: learning from recent behaviour (average and variance)
 - Easily affected by outliers



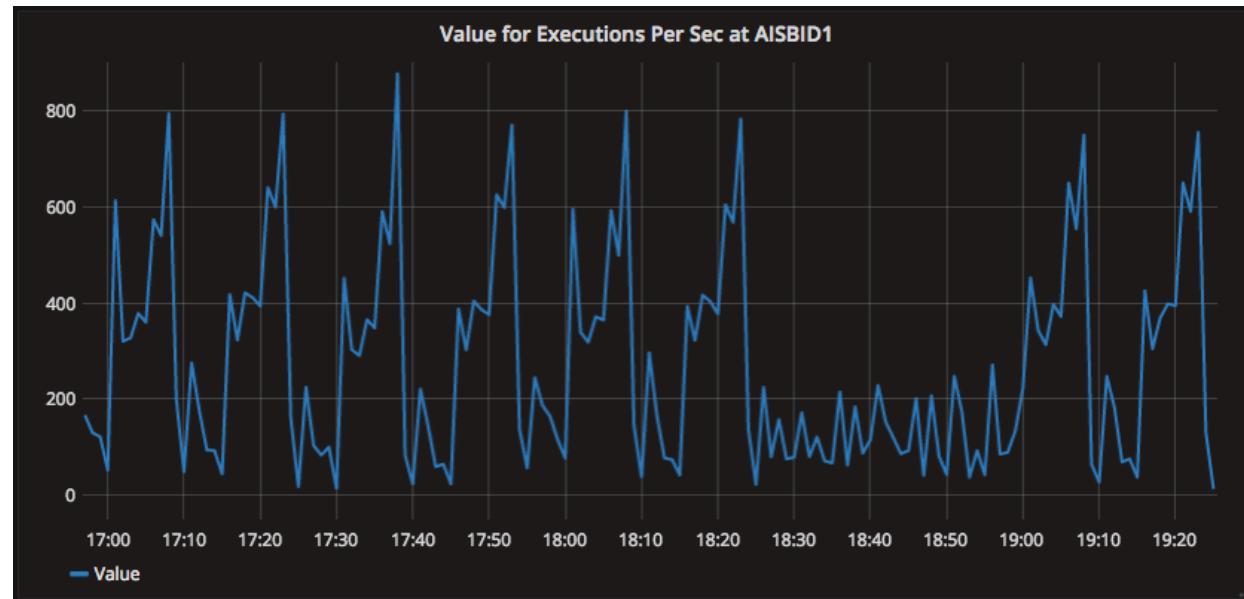


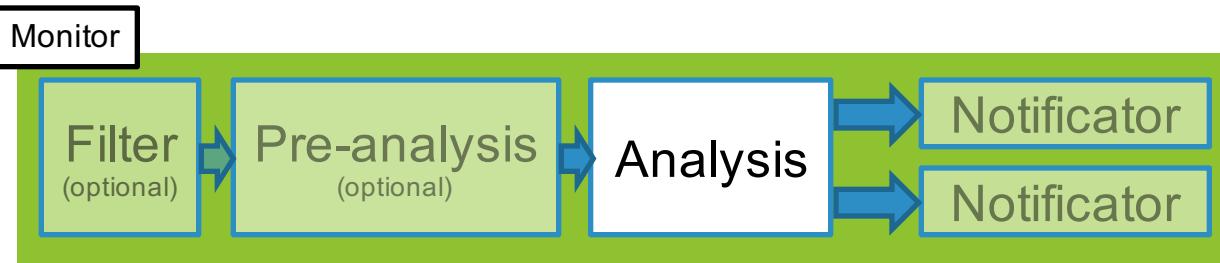
- Type: learning from recent behaviour (percentiles)



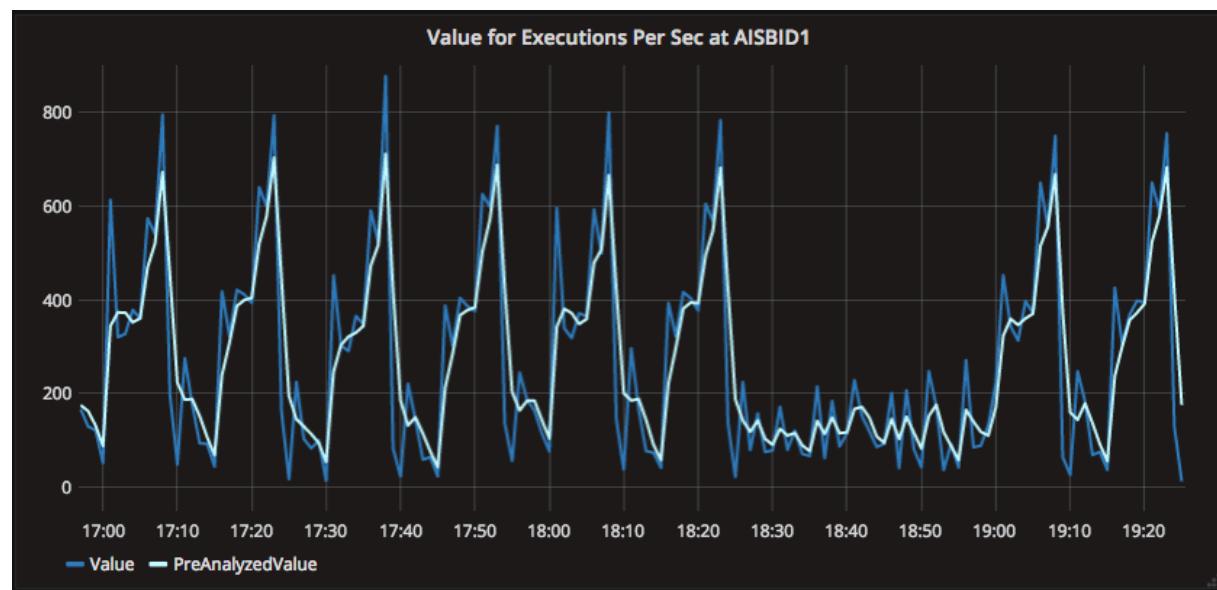


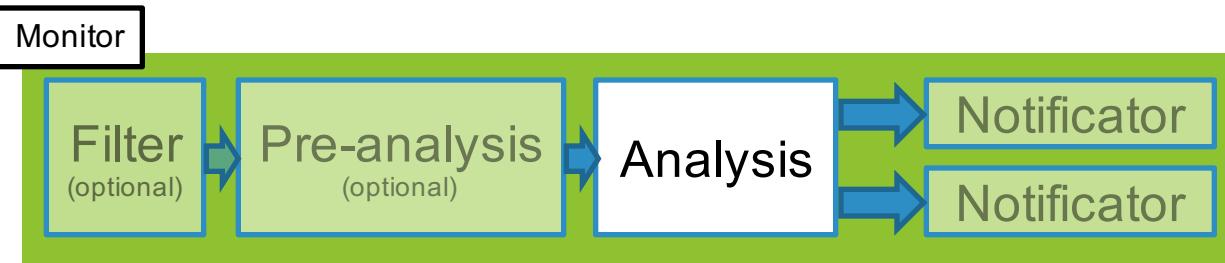
- Type: learning a season (hour, day, week)



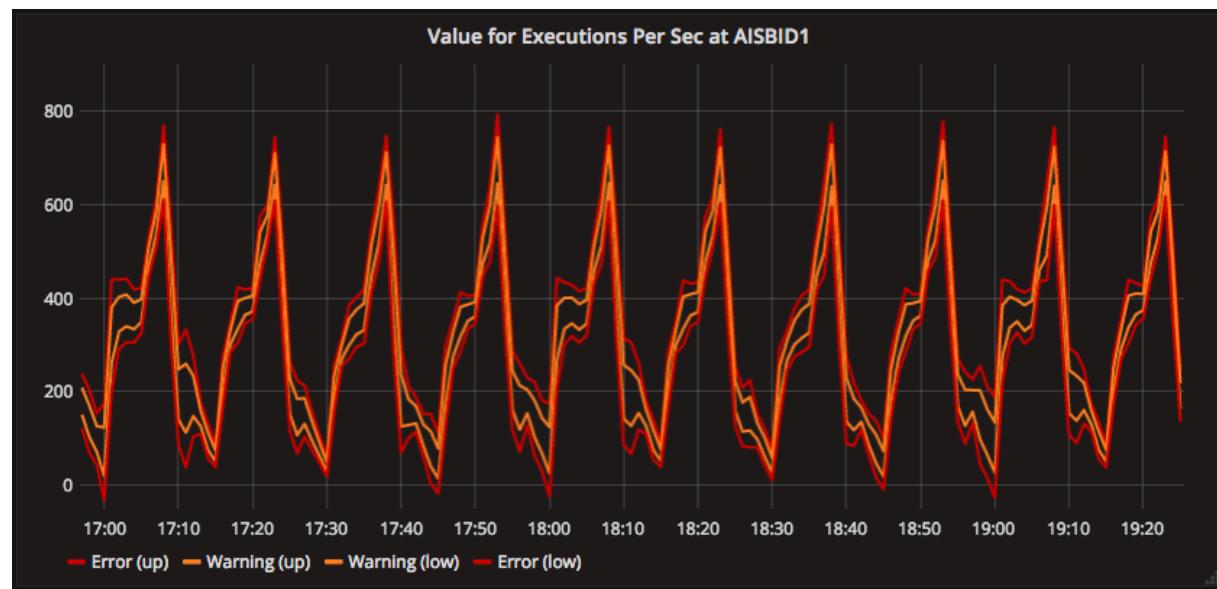


- Type: learning a season (hour, day, week)

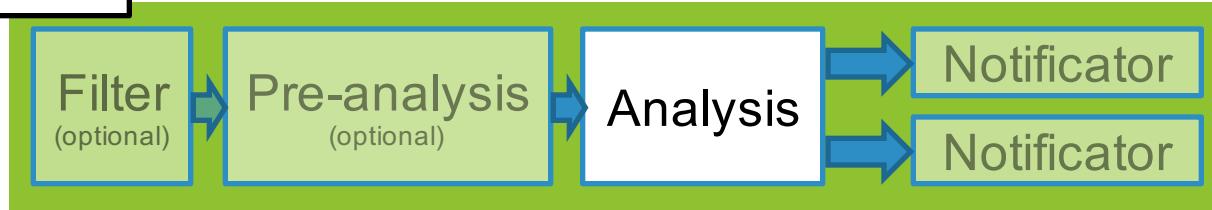




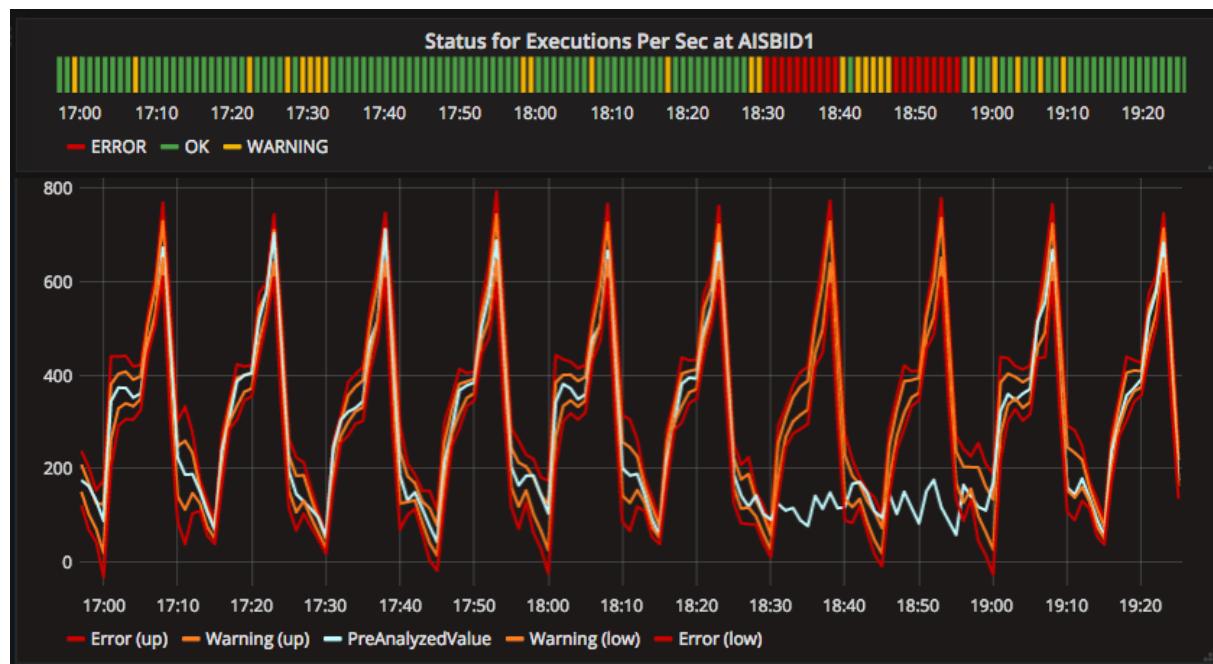
- Type: learning a season (hour, day, week)



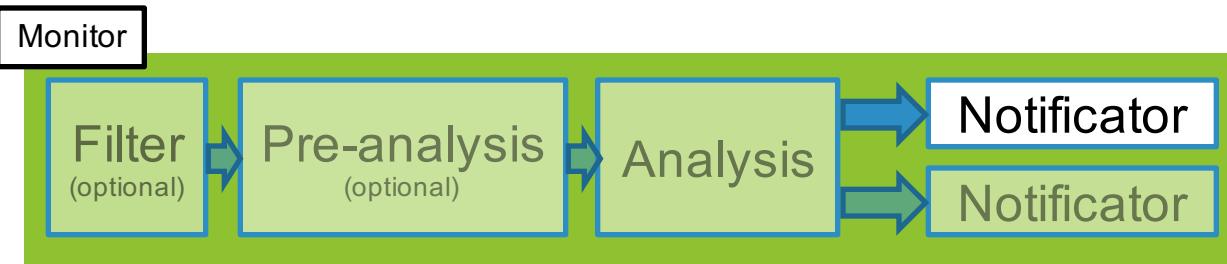
Monitor



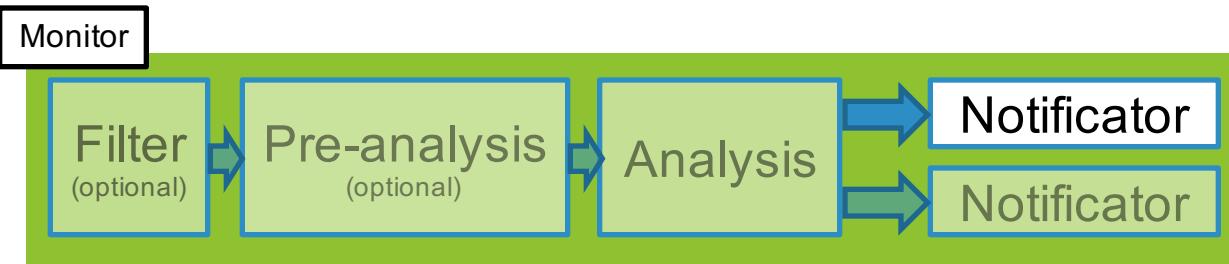
- Type: learning a season (hour, day, week)



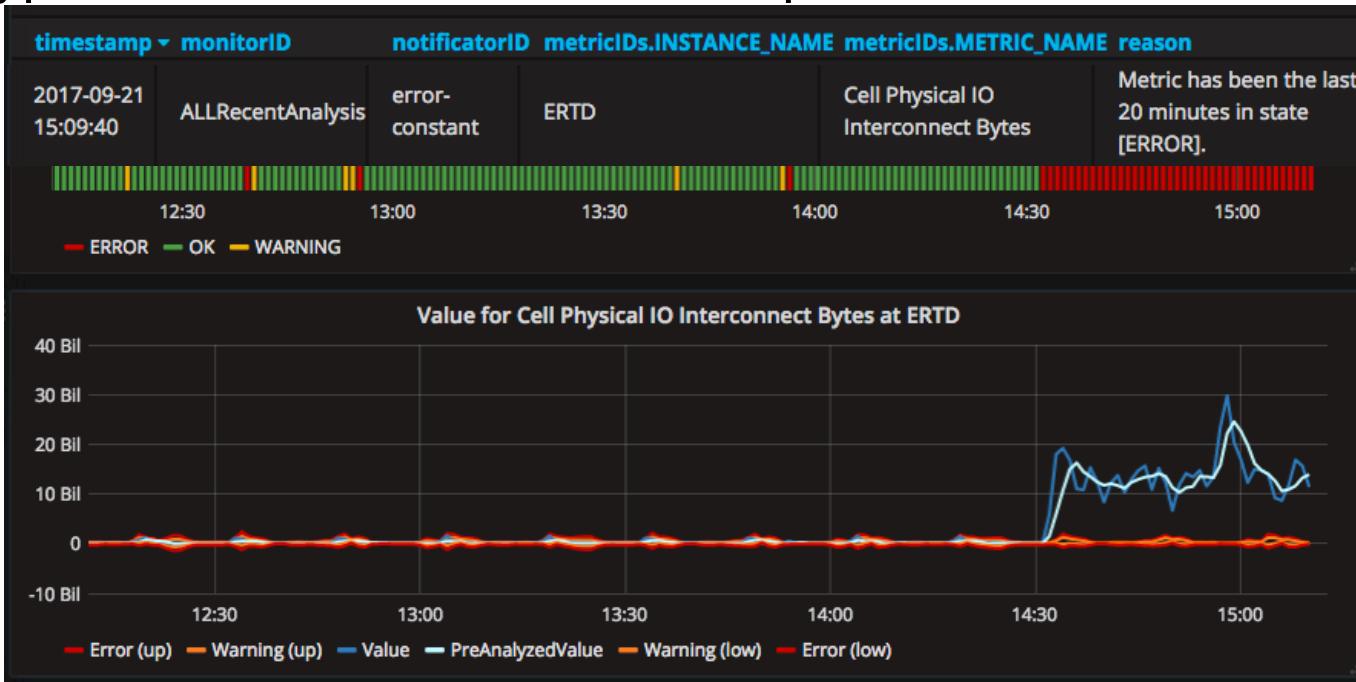
#EUde13



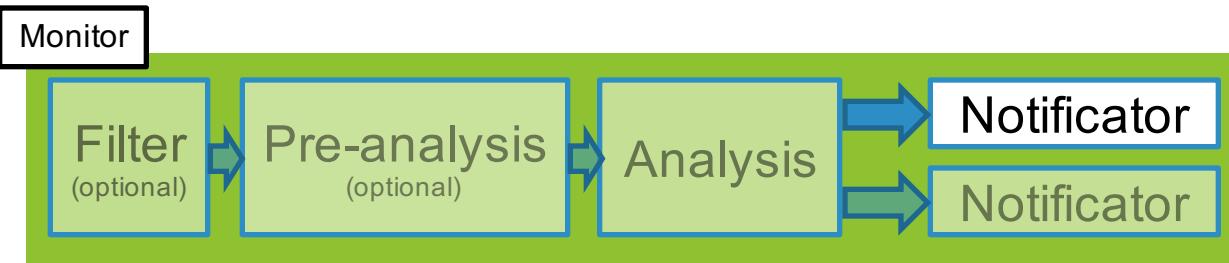
- Should we inform anyone?
 - Metric has been in ERROR for 20 minutes...
- Determines when a notification is raised
- Based on metric statuses
- Notifications are consumed by Notifications sink



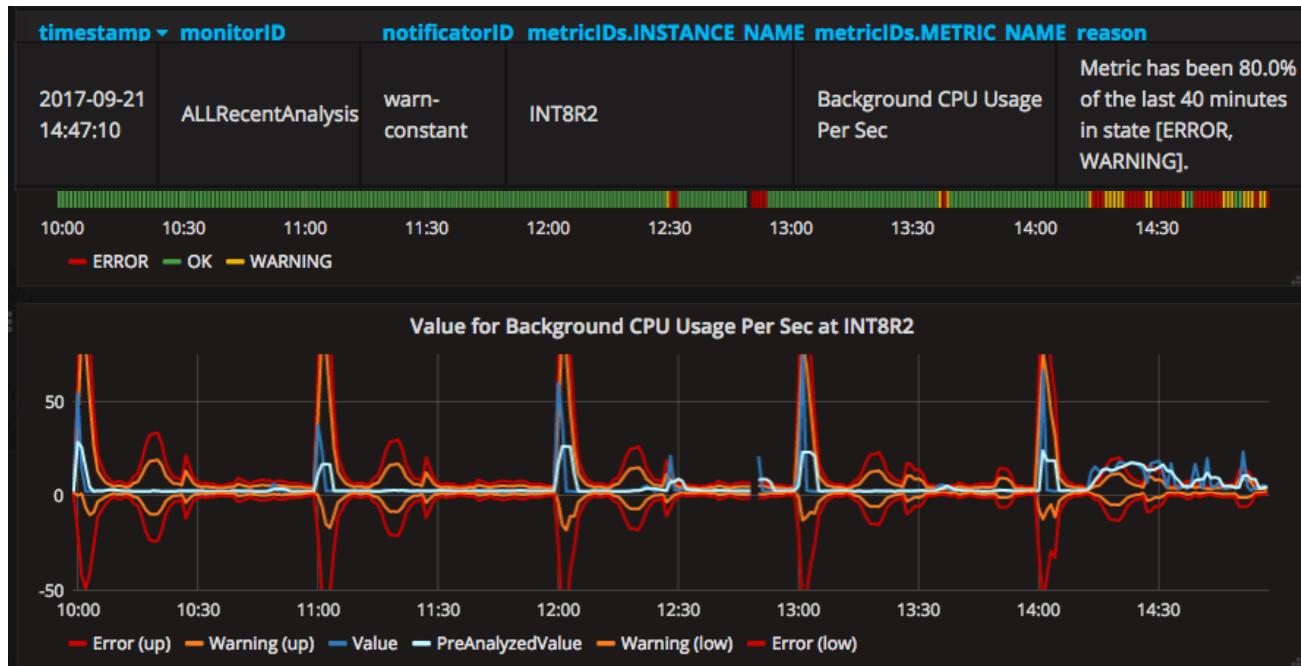
- Type: constant status for certain period



#EUde13



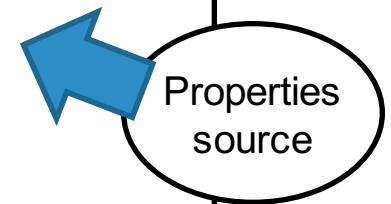
- Type: percentage of the period in certain statuses



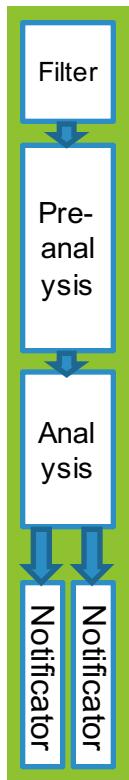
#EUde13

Metrics monitor job: configuration

```
checkpoint.dir = <path_to_store_stateful_data> (default: /tmp/)  
spark.batch.time = <seconds> (default: 30)  
  
metrics.source.kafka-prod.type = kafka  
metrics.source.kafka-prod.consumer.bootstrap.servers = server:port,server:port  
metrics.source.kafka-prod.topics = db-logging-platform  
metrics.source.kafka-prod.parser.attributes = INSTANCE_NAME METRIC_NAME  
metrics.source.kafka-prod.parser.value.attribute = VALUE  
metrics.source.kafka-prod.parser.timestamp.attribute = END_TIME  
  
results.sink.type = elastic  
results.sink.index = itdb_db-metric-results/log  
notifications.sink.type = elastic  
notifications.sink.index = itdb_db-metric-notifications/log  
  
# Monitors  
monitor.<monitor-id-1>.<confs>...  
monitor.<monitor-id-2>.<confs>...  
monitor.<monitor-id-n>.<confs>...
```



Metrics monitor job: monitor configuration



- Configuration can be updated while running
 - Add/remove monitor or change any parameter

```
# Monitor CPU of all instances in production
monitor.CPUUsage.filter.attribute.INSTANCE_NAME = regex:prod-.*
monitor.CPUUsage.filter.attribute.METRIC_NAME = CPU Usage Per Sec
monitor.CPUUsage.missing.max-period = 3m ←
monitor.CPUUsage.pre-analysis.type = weighted-average
monitor.CPUUsage.pre-analysis.period = 5m
monitor.CPUUsage.analysis.type = seasonal
monitor.CPUUsage.analysis.season = hour
monitor.CPUUsage.analysis.learning.ratio = 0.2
monitor.CPUUsage.notifier.error-perc.type = percentage
monitor.CPUUsage.notifier.error-perc.statuses = ERROR
monitor.CPUUsage.notifier.error-perc.period = 10m
monitor.CPUUsage.notifier.warn-perc.type = percentage
monitor.CPUUsage.notifier.warn-perc.statuses = WARNING
monitor.CPUUsage.notifier.warn-perc.period = 20m
```

Metrics monitor job

- We are monitoring:
 - Around 30K metrics per minute (day season analysis)
 - Processing time = 30 seconds/minute
 - On 3 executors (YARN)
- Repo, user's and developer's manual, contributions, issues:
 - <https://github.com/cerndb/metrics-monitor>

Tips and Suggestions

Java 8: `import java.time.*;`

```
private List<DatedValue> values;

private Duration period = Duration.ofMinutes(30);

public void purge(Instant time) {
    Instant oldest_time = time.minus(period);

    values.removeIf(value -> value.getInstant().isBefore(oldest_time));
}
```

Java 8: Extend RDDs and Streams + Optional + lambdas ++ java.streaming.* + reference methods = ❤

```
Stream<Metric> metrics = metricSources.stream()
    .map(source -> source.createStream(ssc))
    .reduce((str, stro) -> str.union(stro)).get();

metrics = metrics.union(metrics.mapS(definedMetrics::generate));

Stream<AnalysisResult> results = metrics.mapS(monitors::analyze);

analysisResultsSink.ifPresent(results::sink);

Stream<Notification> notifications = results.mapS(monitors::notify);

notificationsSink.ifPresent(notifications::sink);
```

Not using Spark check-pointing feature

- **Pro:** develop/update the job without loosing stateful data
- **Cons:** implement save/load

```
JavaRDD<Tuple2<K, S>> initialStates = RDD.<Tuple2<K, S>>.load(getSparkContext(), id);

StateSpec<K, V, S, R> statusSpec = StateSpec
    .function(updateStatusFunction)
    .initialState(initialStates.rdd())
    .timeout(getDataExpirationPeriod());

StatusStream<K, V, S, R> statusStream = StatusStream.from(asJavaPairDStream().mapWithState(statusSpec));

statusStream.getStatuses().save(id);
```

- Note: for stateful operations checkpoint needs to be enable
 - But when restarting, data is removed

Conclusion

- Building the data lake opens up for several kinds of analysis
- Implementation of anomalous connections helps us to identify potential intrusions to the database infrastructure
- Intelligent monitoring tool learns and adapts the threshold
- Intelligent monitoring tool is open source and available for everyone
- Apache Spark Streaming helps us to perform machine learning at scale